

POTENTIALEP: BACKPROP-FREE TRAINING OF CIFAR-SCALE GENERATIVE MODELS VIA EQUILIBRIUM PROPAGATION

Alex Gower*

University of Cambridge & Nokia Bell Labs

ABSTRACT

Generating an image with a diffusion or flow-matching model typically involves integrating a velocity field for hundreds of sequential steps. That computation suits a GPU poorly and an analog neuromorphic substrate well, where one fast, low-power physical relaxation could replace the whole sequence. Two things are needed: a generative algorithm whose computation is a relaxation, and a way to train it without backpropagation, whose global gradient transport is hard to realise in analog. Energy Matching supplies the first, sampling by gradient-descent relaxation on a time-independent potential — but obtaining the gradient that descent follows is itself a backward pass, once per integration step. We prove that a predictive coding network with a single linear injection term added to its energy relaxes to a state whose visible layer carries that gradient inherently, with an error that vanishes for small injection strengths. An Energy Matching model trained by backpropagation can therefore be loaded into this network unchanged and sampled by relaxation alone, matching the original model’s FID to within 0.04 at CIFAR-10 and ImageNet-32 scale. The network can also be trained without backpropagation: Equilibrium Propagation reproduces backpropagation’s parameter updates to a cosine similarity of 0.9999, and training end to end with those updates reaches an FID of 3.64 on unconditional CIFAR-10 against 3.54 for backpropagation. To our knowledge this is the first CIFAR-10 generative model trained with Equilibrium Propagation, and the first at competitive quality by any backpropagation-free rule. This quality does not require group normalisation or attention: an architecture of only convolutions and pointwise nonlinearities, both native to an analog crossbar, reaches an FID of 3.51 when trained by backpropagation.

1 INTRODUCTION

Modern generative models are expensive to train (Strubell et al., 2019) and expensive to sample (Luccioni et al., 2024): diffusion, flow matching and related methods typically produce one image by integrating a velocity field over hundreds of sequential steps, each a full pass through a large network. Sequential work of this kind suits a GPU poorly and a physical system well. A substrate that relaxes to the minimum of an energy performs in one continuous settling what a GPU performs as hundreds of discrete evaluations, and analog neuromorphic devices promise large gains in speed and energy on that basis. Two things stand in the way: finding a generative algorithm whose computation is a relaxation, and removing the backpropagation that training relies on.

Energy Matching (Balcerak et al., 2025) is one such algorithm. It generates by gradient descent on a potential $V_\theta(x)$, the scalar output of a feedforward network, whose minima are where the data lies. The potential does not change with time, so every step descends the same fixed landscape, which is what a physical system does when it relaxes. But the velocity it follows is $-\nabla_x V_\theta$, and in its original formulation obtaining that gradient is a backward pass — once per integration step during sampling, on top of the usual $\nabla_\theta \mathcal{L}$ backward pass per training step.

The remaining obstacle is therefore backpropagation, whose global gradient transport is difficult to realise in analog. Mixed-signal crossbar designs convert between analog and digital at every layer, at

*Correspondence: apg59@cam.ac.uk, or permanently at alexanderpgower@gmail.com

a cost in energy and latency (Kendall et al., 2020). Physics-aware training instead runs the forward pass physically but the backward pass through a digital model of the substrate, and pays for digital computation at every training step (Wright et al., 2022). Both pay a digital cost the physics does not.

The correspondence. We show that the physical dynamics of a predictive coding network can compute $-\nabla_x V_\theta$ inherently. Build a PCN from the same layers, one prediction error per node, and add a single linear term of strength γ on the scalar output node, giving an energy E_{PCN} . The states of this *PotentialPCN* relax to an equilibrium at which $-\nabla_x E_{\text{PCN}} = \gamma(-\nabla_x V_\theta) + O(\gamma^2)$, so dividing by γ recovers the velocity, exactly in the limit of small γ (Section 3.1).

Generation without a backward pass. An Energy Matching model trained conventionally by backpropagation can therefore be loaded into this network unchanged and sampled by relaxation, so that a substrate hosting it needs no backward pass anywhere in the generation path. At CIFAR-10 and ImageNet-32 scale the velocities agree with autograd’s to a cosine similarity of 1 to nine decimal places, and the FID of the generated samples matches to within 0.04 (Section 4.1). No retraining is required, so the energy and latency of generation — most of a deployed model’s lifetime cost — could be cut on analog hardware without any change to how the model is trained.

Training without a backward pass. Training can move on-chip as well. The PCN relaxes on an explicit energy, so it can be trained by Equilibrium Propagation (Scellier & Bengio, 2017), which reads parameter updates $-\nabla_\theta \mathcal{L}$ from local measurements at equilibrium and so needs none of the global transport that makes backpropagation awkward in analog. We use the GradEP variant (Gower, 2026), which trains gradients of energy landscapes — here the velocities Energy Matching generates with. We call the resulting procedure *PotentialEP*. Its parameter updates reproduce backpropagation’s to a cosine similarity of 0.9999 at our operating point (Section 4.2), and a model trained end to end with those updates reaches an FID of 3.64 on unconditional CIFAR-10, against 3.54 for backpropagation (Section 5). To our knowledge this is the first CIFAR-10 generative model trained with Equilibrium Propagation, and the first at competitive quality by any backpropagation-free rule.

The architecture. A backpropagation-free algorithm is of little use if the architecture it trains is itself hard to realise in analog, and two components of the one we inherit from Energy Matching are: group normalisation, which couples every channel through a shared statistic, and attention, whose pairwise products are not a fixed weighted sum. We find that removing both costs nothing measurable at this scale — a network of only convolutions, pointwise nonlinearities and UNet skips, every operation of the kind an analog crossbar performs natively, reaches a CIFAR-10 FID of 3.51 (Section 5.1).

2 BACKGROUND

2.1 ENERGY MATCHING

Energy Matching (Balcerak et al., 2025) generates images by relaxation on a learned scalar potential. A feedforward network with a scalar output defines

$$V_\theta(x) = \alpha F_\theta(x), \quad (1)$$

where $F_\theta(x)$ is the network output and α a fixed output scale set by the architecture, and samples are produced by integrating $\dot{x} = -\nabla_x V_\theta(x)$ from noise. Unlike diffusion or flow matching, the velocity field is the gradient of a *time-independent* potential: the same landscape is descended at every step, and the model’s minima are where the data lies.

Training has two phases. The first is flow matching (Lipman et al., 2023) on exact optimal-transport couplings (Tong et al., 2024): for a noise sample x_0 paired with a data sample x_1 , an interpolant $x_t = (1-t)x_0 + tx_1$ with $t \sim \mathcal{U}[0, 1]$, and a target velocity $u = x_1 - x_0$, the loss is

$$\mathcal{L}(\theta) = \frac{1}{2} \| -\nabla_x V_\theta(x_t) - u \|^2. \quad (2)$$

The coupling is needed because the field $-\nabla_x V_\theta$ is time-independent and curl-free by construction, while the marginal velocity it is asked to match need be neither; pairing each x_0 with its transport-optimal x_1 reduces the overlap between conditional trajectories, so that the marginal they define is

close to time-independent and curl-free. The second phase adds a contrastive term, $\nabla_{\theta} V(x^+) - \nabla_{\theta} V(x^-)$ between data samples and model samples produced by a short Langevin chain. We refer to these throughout as phases 1 and 2.

Sampling integrates

$$dx = -\nabla_x V_{\theta}(x) d\tau + \sqrt{\epsilon(\tau)} dW \quad (3)$$

to a stopping time τ_s , with the noise ϵ switched off below $\tau = 1$ and held at a fixed level beyond it: up to $\tau = 1$ it is a noiseless ODE transporting noise to data, and beyond that a Langevin chain at fixed temperature. Because the landscape does not change with time, integrating past $\tau = 1$ is further equilibration rather than extrapolation into an untrained regime, as it would be for a time-conditioned flow-matching model, so τ_s is a free parameter, chosen per model over a ladder of values (App. J).

2.2 PREDICTIVE CODING NETWORKS

A predictive coding network (Rao & Ballard, 1999; Whittington & Bogacz, 2017) assigns a dynamical variable h_k to each node of a feedforward architecture and an energy to the mismatch between each node and the prediction made of it by its predecessor:

$$E = \sum_{k=1}^N \frac{1}{2} \|f_k(h_{k-1}) - h_k\|^2, \quad e_k = f_k(h_{k-1}) - h_k. \quad (4)$$

Inference is not a forward pass but a relaxation: with the input $h_0 = x$ clamped, the states descend E to a stationary point, and outputs are read from the settled state. Two properties of that stationary state are what Section 3.1 builds on. First, with only error terms present the minimum is $e = 0$, which by induction from $h_0 = x$ is exactly the forward pass: left to itself the relaxation reproduces the feedforward network it was built from. Second, stationarity $\partial E / \partial h_k = 0$ at an interior node gives $e_k = J_{k+1}^{\top} e_{k+1}$ with $J_{k+1} = \partial f_{k+1} / \partial h_k$, so any error introduced at the top propagates downwards through transposed Jacobians.

Each term of equation 4 involves only a node and its predecessor, so both the relaxation and the parameter gradients $\partial E / \partial \theta_k = e_k^{\top} \partial f_k / \partial \theta_k$ are local: every quantity a weight needs is present at its two ends. This is what makes predictive coding a candidate for analog implementation, where the relaxation is performed by the physics rather than simulated.

2.3 EQUILIBRIUM PROPAGATION

Equilibrium Propagation (Scellier & Bengio, 2017; Laborieux et al., 2020) is an alternative learning algorithm to backpropagation for systems whose dynamics settle to stationary points, from which their outputs are read. It uses only local measurements at equilibrium, making it a natural candidate for neuromorphic hardware.

Consider a system with state s , parameters θ and energy $E(s; \theta)$, relaxing to an equilibrium s^* whose position we want to train so that the outputs reduce a cost $C(s)$. Adding $\beta C(s)$ to the energy displaces that equilibrium towards lower cost, and the parameter gradient can be read from where and how far it moves. Each training sample therefore takes three phases: a *free phase* relaxing to s^* , then two *nudged phases* relaxing $E \pm \beta C$ to $s^{*\pm\beta}$, at lower and higher cost respectively. The update is

$$\Delta\theta \propto -\frac{1}{2\beta} \left[\frac{\partial E}{\partial \theta} \Big|_{s^{*+\beta}} - \frac{\partial E}{\partial \theta} \Big|_{s^{*-\beta}} \right] \xrightarrow{\beta \rightarrow 0} -\nabla_{\theta} C(s^*), \quad (5)$$

which lowers the energy at $s^{*+\beta}$ and raises it at $s^{*-\beta}$, so that the free equilibrium s^* itself shifts towards lower cost. For small enough β this is exactly gradient descent on $C(s^*)$.

Each $\partial E / \partial \theta$ is a local quantity — on a PCN, $e_k^{\top} \partial f_k / \partial \theta_k$ — so the whole procedure runs without global backward passes. EP has been scaled to CIFAR-10 and ImageNet classification, including on PCN energies (Kerjan et al., 2026), but its generative applications remain at MNIST scale and below (Meersch et al., 2023; Gower, 2026).

The cost C must be a function of s^* itself, the position the state settles at, since that is what the nudge displaces. A loss involving an energy *gradient* at that position — such as Energy Matching’s

equation 2 involving $\nabla_x V$ — is therefore out of reach. GradEP (Gower, 2026) resolves this by running EP on a modified energy whose equilibrium position encodes the gradient of the original energy, so that moving the former’s equilibrium trains the latter’s gradient. Section 3.2 applies this to the PCN of Section 3.1, and Section 3.2.1 improves the nudge it uses.

3 METHOD

3.1 A PCN WHOSE RELAXATION COMPUTES $-\nabla_x V$

Energy Matching defines its potential as the scalar output of a feedforward network, scaled by a fixed constant α : $V_\theta(x) = \alpha F_\theta(x)$, with $F_\theta = f_N \circ \dots \circ f_1$ the network and $F_\theta(x)$ its scalar output. Generation integrates $\dot{x} = -\nabla_x V$, and every evaluation of that gradient is a backward pass through F_θ . We now build, from the same layers f_k , a predictive coding network that produces the same gradient as a property of its relaxed state, reached by dynamics that are local at every node.

Take the standard PCN energy for this network — one prediction error per node — and add a single linear *injection* term on the output node:

$$E_{\text{PCN}}(x, h) = \sum_{k=1}^N \frac{1}{2} \|f_k(h_{k-1}) - h_k\|^2 + \gamma \alpha h_N, \quad h_0 = x, \quad h_N = o \in \mathbb{R}, \quad (6)$$

where $h = (h_1, \dots, h_N)$ are the node states, $h_N = o$ is the scalar output node (α is kept separate from γ so that $\gamma > 0$ is expressed on the potential’s scale), and $e_k = f_k(h_{k-1}) - h_k$ are the errors. We call this construction a *PotentialPCN*. The hidden states relax at fixed x to a stationary point $h^*(x)$ of E_{PCN} , and we read a velocity from the visible layer:

$$v_{\text{PCN}}(x) = -\gamma^{-1} \nabla_x E_{\text{PCN}}(x, h^*(x)). \quad (7)$$

Proposition 1. *For smooth f_k and $\gamma \rightarrow 0$,*

$$v_{\text{PCN}}(x) = -\nabla_x V_\theta(x) + O(\gamma), \quad \gamma^{-1} E_{\text{PCN}}(x, h^*(x)) = V_\theta(x) + O(\gamma^2). \quad (8)$$

The argument is short, and App. A gives it in full. Without the injection the energy is minimised at $e = 0$, which by induction from $h_0 = x$ is the forward pass. The injection disturbs this: stationarity in the output node gives $e_N^* = \gamma \alpha$ exactly, and stationarity in an interior node gives $e_k^* = J_{k+1}^\top e_{k+1}^*$ with $J_{k+1} = \partial f_{k+1} / \partial h_k$, so every error is $O(\gamma)$ and each node h_k^* sits $O(\gamma)$ from the forward-pass state it would occupy at $\gamma = 0$. The velocity is the total derivative of $E_{\text{PCN}}(x, h^*(x))$ in x , which has an explicit part and a part through $h^*(x)$; the latter vanishes because h^* is stationary in h , and iterating the interior recursion from the top leaves $\nabla_x E_{\text{PCN}}(x, h^*(x)) = J_1^\top e_1^* = \gamma \alpha J_1^\top J_2^\top \dots J_N^\top$. Evaluating each Jacobian at the forward pass rather than at h^* costs $O(\gamma)$, and the product is then the chain rule for $\nabla_x F_\theta$, so $\nabla_x E_{\text{PCN}}(x, h^*(x)) = \gamma \alpha \nabla_x F_\theta(x) + O(\gamma^2)$. For the second claim, evaluate E_{PCN} at h^* : the squared errors are $O(\gamma^2)$, so the injection term dominates, giving $\gamma V_\theta(x) + O(\gamma^2)$.

The PCN’s errors *are* the backward pass: the linear term injects the constant $\gamma \alpha$ at the output, analogous to the “ $\partial V / \partial V = 1$ ” that autograd seeds its backward pass with, and the relaxation propagates it down through transposed Jacobians, exactly as backpropagation would. The velocity is then read from the first node alone, $-\gamma^{-1} J_1^\top e_1^*$: where backpropagation must traverse all N layers to reach x , here that traversal has already happened in the relaxation, and the visible layer sees only its own local error.

Nothing in the derivation used the values of the weights: it holds for any θ , so a network trained conventionally as F_θ can be loaded into equation 6 unchanged and sampled by relaxation (Section 4.1). What the derivation does assume is that h is at equilibrium for the current x . In simulation we enforce this by alternating: x takes an integration step, then h is relaxed by K_h sweeps of gradient descent on E_{PCN} at step dt (which in the error parameterisation of App. E converges in a single sweep $K_h = 1$ at $dt = 1$). A physical realisation needs no alternation: if the hidden nodes relax on a faster timescale than x , h tracks $h^*(x)$ adiabatically and the generation trajectory is one continuous relaxation. The generalisation from a chain to the UNet’s skip-connected graph replaces $J_{k+1}^\top e_{k+1}$ by a sum over children and is given in App. D.

3.2 TRAINING THE PCN BY EQUILIBRIUM PROPAGATION

Section 3.1 shows that the PCN’s relaxation reproduces the generation velocity $-\nabla_x V_\theta$ for any weights θ , whether or not those weights were themselves trained on the PCN. Training them on the PCN requires two more things: that the correspondence hold for the loss’s *parameter* gradient and not only for the velocity, and a way to measure that parameter gradient without a backward pass.

The PCN’s loss is the potential’s loss. In its first phase Energy Matching trains V_θ by flow matching: for an interpolant x_t between a noise sample and a data sample, with target velocity u , the loss is $\mathcal{L}(\theta) = \frac{1}{2} \|\nabla_x V_\theta(x_t) - u\|^2$, and $\Delta\theta_{\text{BP}} = -\nabla_\theta \mathcal{L}$ is its backpropagated gradient. To allow this loss (dependent on an energy *gradient*) to be trained using EP, we use GradEP’s spring-clamped free phase: rather than fixing x at x_t , we tether x near x_t with a spring energy term and let x and h relax together on

$$E_{\text{free}}(x, h) = E_{\text{PCN}}(x, h) + \frac{\lambda}{2} \|x - x_t\|^2. \quad (9)$$

Stationarity in x gives $\nabla_x E_{\text{PCN}} = -\lambda(x^* - x_t)$, so by equation 7 the velocity we want to train is encoded in the spring’s equilibrium extension,

$$v^* = -\gamma^{-1} \nabla_x E_{\text{PCN}}(x^*, h^*) = \frac{\lambda}{\gamma} (x^* - x_t), \quad (10)$$

and the PCN’s loss is $\mathcal{L}_{\text{PCN}}(\theta) = \frac{1}{2} \|v^* - u\|^2$. Training is now a matter of moving where x settles on this spring-clamped energy landscape: the weights must be adjusted so that the spring’s equilibrium extension encodes a velocity closer to u , which is the kind of objective EP can act on. The correspondence error varies smoothly with θ as well as being small at each (App. A), so $\nabla_\theta \mathcal{L}_{\text{PCN}} = \nabla_\theta \mathcal{L} + O(\gamma)$ and descending $\mathcal{L}_{\text{PCN}}$ descends the potential’s own loss: EP shapes the E_{PCN} landscape and V_θ is trained as a consequence, a claim Section 4.2 tests over six decades of γ .

Two ways to compute the gradient. Because h^* is defined implicitly by stationarity, the parameter gradient of the flow loss on $v_{\text{PCN}}(x_t)$, with x held fixed at x_t rather than tethered, can be obtained by implicit differentiation, solving a linear system in the relaxation Hessian (App. E). We write this update $\Delta\theta_{\text{IFT}}$ and use it as the control in Section 4.2: it carries none of EP’s finite-difference error, only the $O(\gamma)$ correspondence bias (App. H). Its matrix–vector products use autograd, however, so a neuromorphic substrate would run Equilibrium Propagation instead.

The three phases of EP. GradEP realises equation 5 in three phases, using the centred $\pm\beta$ form of Laborieux et al. (2020): a free phase that relaxes equation 9 and reads v^* , then two nudged phases that add $\pm\beta' C(x)$ to it and relax again from the free equilibrium. The cost is the loss as a function of the visible state, $C(x) = \frac{1}{2} \|v(x) - u\|^2$ with $v(x) = \frac{\lambda}{\gamma} (x - x_t)$, so that $C(x^*) = \mathcal{L}_{\text{PCN}}$. Its Hessian in x is $(\lambda/\gamma)^2 I$, inherited from the velocity conversion, so we nudge with $\beta' = \beta (\gamma/\lambda)^2$: the nudge then contributes βI against the tether’s λI , making β a stiffness directly comparable to λ . Equation 5 then gives

$$\Delta\theta_{\text{EP}} = -\frac{1}{2\beta'} \left[\frac{\partial E_{\text{PCN}}}{\partial \theta}(s^{*+\beta}) - \frac{\partial E_{\text{PCN}}}{\partial \theta}(s^{*-\beta}) \right] = -\nabla_\theta \mathcal{L}_{\text{PCN}} + O(\beta^2), \quad (11)$$

the second equality using $\nabla_\theta C(x^*) = \nabla_\theta \mathcal{L}_{\text{PCN}}$, and the size of the $O(\beta^2)$ term depending on the choice of cost in a way Section 3.2.1 addresses. Only E_{PCN} appears in equation 11: the tether and the nudge are θ -free. App. B gives the sweep budgets, the choice of λ , how $\mathcal{L}_{\text{PCN}}$ differs from $\mathcal{L}$ at finite γ , and the contrastive phase, which on the PCN needs no nudge at all.

3.2.1 THE LINEARISED NUDGE

GradEP’s cost is the loss as a function of the visible state, $C(x) = \frac{1}{2} \|v(x) - u\|^2$, and being quadratic in x , its nudge acts as a second spring (of strength β) against the tether (of strength λ). Balancing the two forces gives a nudged equilibrium $x^{*\beta} - x^* = -\beta\gamma g / \lambda(\lambda + \beta)$, with $g = v^* - u$ (App. C), which is nonlinear in β . The $\pm\beta$ displacements are therefore asymmetric about the free equilibrium, and the centred difference equation 11 acquires an $O((\beta/\lambda)^2)$ bias. This is standard EP’s weak-nudge requirement, with the tether’s λ in place of an ordinary network’s $O(1)$ curvature. At $\lambda = 1$ it caps β near one, which is restrictive since the displacement the parameter update is read from scales with $\beta\gamma$ and must stay resolvable in finite-precision arithmetic, while γ must stay small for the correspondence.

We therefore instead nudge with the loss’s *linearisation* at the free equilibrium, as in Dual Propagation (Høier et al., 2023): freezing g after the free phase and taking $C_{\text{lin}}(x) = g^\top v(x)$, which targets the same parameter gradient but adds no curvature in x . Its gradient $\nabla_x(\beta' C_{\text{lin}}) = \beta\gamma g/\lambda$ is a constant force, where the quadratic cost’s depends on x and so makes the equilibrium nonlinear in β . The displacement $x^{*\beta} - x^* = -\beta\gamma g/\lambda^2$ is therefore linear in β at any nudge strength and the $\pm\beta$ equilibria stay symmetric about the free one. What remains is only that $\partial_\theta E_{\text{PCN}}$ varies nonlinearly along the straight path between them, so the estimate at finite β departs from the $\beta \rightarrow 0$ limit it targets by a relative error of order $|x^{*\beta} - x^*|^2$ (App. C), and the constraint becomes

$$\beta\gamma|g| \ll \lambda^2 \quad \text{in place of} \quad \beta \ll \lambda, \quad (12)$$

so smaller γ permits larger β : at the $\gamma = 10^{-4}$, $\lambda = 1$ and $|g| \approx 0.39$ of our phase-1 training runs, the ceiling moves from $\beta \ll 1$ to $\beta \ll 2.6 \times 10^4$. Section 4.2 measures the resulting window in $\beta\gamma$.

4 EQUIVALENCE EXPERIMENTS

4.1 EQUIVALENCE AT INFERENCE

We first compare, at identical trained weights, the velocities obtained by automatic differentiation through the feedforward Energy Matching model, $v_{\text{FFN}}(x) = -\nabla_x V(x)$, against those produced by the relaxed PCN built from it (App. D), $v_{\text{PCN}}(x) = -\gamma^{-1}\nabla_x E_{\text{PCN}}(x, h)|_{h^*}$. This tests the correspondence of Section 3.1 numerically.

The inference velocity correspondence is $O(\gamma)$ down to machine precision. Figure 1 displays the relative velocity deviation $\epsilon = \|v_{\text{PCN}} - v_{\text{FFN}}\|/\|v_{\text{FFN}}\|$ as a function of γ , for our replication of the CIFAR-10 Energy Matching model (Balcerak et al., 2025). The deviation decays proportional to γ over six decades, from $\gamma = 0.1$ down to $\gamma = 10^{-7}$, in both its direction and magnitude components (App. F), as predicted by the derivation in App. A. At small γ each arithmetic saturates at its own precision floor (App. F); the highest is the $\sim 2 \times 10^{-4}$ of the TF32 tensor-core mode our runs use. Above $\gamma \sim 10^{-2}$ the curves coincide, the finite- γ bias having grown past every floor. At the $\gamma = 10^{-4}$ of our CIFAR-10 runs the float64 total relative deviation is 1.3×10^{-5} (a direction cosine similarity of $\cos(v_{\text{FFN}}, v_{\text{PCN}}) = 1$ to nine decimal places), which lies below the TF32 floor. The correspondence error for our runs is therefore dominated by standard GPU arithmetic precision rather than by the finite- γ bias of the relaxation, so smaller γ buys no accuracy in this arithmetic. Repeating these measurements on the authors’ released ImageNet-32 checkpoints gives the same behaviour (App. G), so the correspondence is not specific to this dataset or trained potential.

Generation composes hundreds of velocity evaluations, so we also compare the two velocity fields end to end by sampling with each and measuring the FID of the generated samples.

Generation by relaxation. We run the Energy Matching sampler unchanged, replacing v_{FFN} with v_{PCN} at every integration step, on our CIFAR-10 replication checkpoints and on the authors’ released ImageNet-32 checkpoints — a scale where training is beyond our compute budget but sampling is not. Table 1 shows the FID matches feedforward inference in every case. Within each row both runs start from identical noise, so the differences reflect the velocity source rather than sampling variance; the largest is 0.036; App. F shows samples from both velocity sources side by side. A trained Energy Matching model can therefore be sampled at ImageNet-32 scale by relaxation alone, with no measurable quality loss.

4.2 EQUIVALENCE AT TRAINING

We now ask what training on the PCN requires, comparing the updates it produces — by EP, $\Delta\theta_{\text{EP}}$, and by the implicit-function solve of Section 3.2, $\Delta\theta_{\text{IFT}}$ — against the backpropagation gradient $\Delta\theta_{\text{BP}}$ of the same feedforward model on the same batch. The optimiser rescales the update magnitude, so we measure only its direction, as a cosine similarity with $\Delta\theta_{\text{BP}}$. We relax the free phase to convergence throughout (App. H).

The correspondence itself is not the constraint. Where EP estimates the parameter gradient by a finite difference of two nudged equilibria, the implicit solve obtains it directly at equilibrium, so




Figure 1: Relative deviation $\epsilon = \|v_{\text{PCN}} - v_{\text{FFN}}\| / \|v_{\text{FFN}}\|$ of the relaxed PCN velocity from the exact FFN autograd velocity, on the CIFAR-10 replication checkpoint of Table 1, evaluated at 512 flow-matching interpolants x_t constructed from training images exactly as during training. The deviation is $O(\gamma)$ (dashed) until it meets the arithmetic floor of the number format; float64 meets no floor in the measured range, and at the operating $\gamma = 10^{-4}$ (dotted) the deviation lies below the TF32 floor of our runs.

Table 1: Backprop-trained checkpoints sampled with feedforward autograd velocities (v_{FFN}) versus PCN relaxation velocities (v_{PCN}), at $\gamma = 10^{-4}$ (CIFAR-10) or 10^{-6} (ImageNet-32). Phases 1 and 2 are the Energy Matching training stages of Section 2; sampling seeds are matched within each row.

Dataset	Checkpoint	FID ↓	
		Feedforward	PCN relaxation
CIFAR-10	Phase 1, $\tau_s = 1.0$	6.42	6.42
	Phase 2, $\tau_s = 3.25$	3.54	3.55
ImageNet-32	Phase 1, $\tau_s = 1.0$	8.33	8.33
	Phase 2, $\tau_s = 2.5$	6.75	6.78

its only error is the $O(\gamma)$ correspondence bias of Section 4.1. Even in float32 $\cos(\Delta\theta_{\text{IFT}}, \Delta\theta_{\text{BP}})$ is exact to four decimal places from $\gamma = 10^{-7}$ to 10^{-2} , softening only to 0.998 at $\gamma = 10^{-1}$, and is equally insensitive to the number of conjugate-gradient steps (App. H). Everything that follows is therefore a test of EP’s finite-difference estimator rather than of the correspondence.

γ and β set a window together. Once the free phase has converged, the fidelity of $\Delta\theta_{\text{EP}}$ to $\Delta\theta_{\text{BP}}$ depends on the size of the displacement the nudge produces, which γ and β control jointly: the two nudged equilibria differ by $O(\gamma\beta)$ (Section 3.2.1), so it is their product that matters. Table 3 (App. H) gives $\cos(\Delta\theta_{\text{EP}}, \Delta\theta_{\text{BP}})$ on a γ - β grid, and the usable region is a diagonal band, where each decade lost in γ is recovered by roughly a decade gained in β : $\gamma = 3 \times 10^{-5}$ needs $\beta \geq 10$, where $\gamma = 10^{-3}$ is clean from $\beta = 0.3$. The two edges have different causes. Below the band the displacement is too small to survive float32, and float64 rescues it completely — 1.0000 against float32’s 0.51 at the same β and $\gamma = 10^{-5}$. Above it the nudged equilibria are displaced too far for a linear finite difference to estimate the derivative well; float64 does not rescue that edge, and since the implicit solve carries the same $O(\gamma)$ bias yet degrades only to 0.998 there, the failure belongs to the nudge rather than to the correspondence or the arithmetic. The band therefore closes at both ends, with no β reaching 0.99 by $\gamma = 10^{-1}$; our phase-1 operating point, $\gamma = 10^{-4}$ with $\beta = 30$, sits comfortably inside it at 0.9999. The window width is due to the linearised nudge: sweeping β alone at the operating γ (App. H), the linear nudge holds $\cos \geq 0.99$ from $\beta = 2$ to 10^3 , where the quadratic nudge only does so at $\beta = 1$ and nowhere above.

Table 2: Unconditional CIFAR-10 FID, with our models after both training phases at each checkpoint’s optimal τ_s . The upper block is taken from the cited papers; * marks reproductions by Balcerak et al. (2025). [‡]Phase-2 negatives drawn with a feedforward twin (see text). Protocol and per-arm differences: App. J, App. I.

Method	Gradient	Params	FID ↓
ImprovedCD (Du et al., 2021)	backprop	—	25.1
SDiT (spiking) (Yang et al., 2026)	backprop	77M	11.34
CLEL-large (Lee et al., 2022)	backprop	31M	8.61
OT-CFM* (Tong et al., 2024)	backprop	37M	4.04
DDPM++ (Kim et al., 2022)	backprop	62M	3.45
Energy Matching (Balcerak et al., 2025)	backprop	50M	3.34
Energy Matching (our replication)	backprop	50M	3.54 ($\tau_s = 3.25$)
PotentialEP	implicit solve	50M	3.64 ($\tau_s = 3.25$) [‡]
PotentialEP	EP	50M	3.64 ($\tau_s = 3.25$) [‡]
Energy Matching, norm- and attention-free (§5.1)	backprop	78M	3.51 ($\tau_s = 5.0$)

5 CIFAR-10 GENERATION BY BACKPROP-FREE TRAINING

The previous section established both correspondences at fixed weights. We now train with them end to end, over 10^5 steps, and measure the generation quality of the resulting model.

We train three arms: the feedforward Energy Matching model by backpropagation, and the PCN at $\gamma = 10^{-4}$ by the implicit solve and by EP with the linearised nudge. The architecture, optimiser, schedule, batch size, EMA decay and step budget are shared; initialisation and the random draws of training are not. The EP arm’s phase-2 differs further: a smaller γ , a stiffer tether λ and a larger β , starting from its 100k phase-1 checkpoint rather than 145k. Relaxation settings are in Table 6 and the rest of the recipe in App. J.

Backprop-free training matches backpropagation. Table 2 gives the three arms. After phase 1, backpropagation reaches an FID of 6.42 at 145k steps and the implicit solve 6.70, while the EP arm reaches 6.23 at the 100k checkpoint its contrastive phase starts from. These span 0.47, more than the ~ 0.2 below which we cannot resolve two FID evaluations (App. J), which is unsurprising for runs that share neither an initialisation nor a data realisation. The contrastive phase then brings all three together: 3.54 for backpropagation, 3.64 for the implicit solve, and 3.64 for EP — the one estimator of the three that a neuromorphic substrate could run. The PCN is therefore trainable end to end without backpropagation in the parameter updates of either phase. Both PCN arms do draw their phase-2 negatives with a feedforward twin carrying the same weights, since relaxing them was beyond our compute budget; no parameter gradient passes through that twin, and Section 4.1 shows the two velocity fields agree (App. J).

The upper block of Table 2 places these numbers among conventionally trained models. Our replication reaches 3.54 against the published 3.34, a difference comparable to the repeat spread of the measurement itself (App. J); on that scale the backprop-free arms sit within the same band as recent diffusion and flow-matching models, and well ahead of prior energy-based ones. To our knowledge this is the first CIFAR-10 generation result for a model trained by Equilibrium Propagation.

5.1 HARDWARE-PLAUSIBLE ARCHITECTURES

PotentialEP is neuromorphically compatible as an algorithm, but the energy it trains is only as neuromorphically implementable as the architecture defining it. Two components we inherit from Energy Matching are not easily realised on an analog substrate: group normalisation, which couples every channel at a layer through a shared statistic, and attention, whose pairwise products are not a fixed weighted sum. We therefore ask how much generation quality is lost when both are removed.

Replacing them with weight standardisation, which normalises each filter’s weights rather than its activations, residual scaling by $1/\sqrt{\text{depth}}$, and a flat linear head in place of the ViT (App. I) leaves

an architecture with no operation beyond fixed weighted sums and pointwise nonlinearities — both native to analog crossbar hardware, the sums through Ohm’s and Kirchhoff’s laws (Ielmini & Wong, 2018; Sebastian et al., 2020). We train this arm by backpropagation, our compute budget not extending to a second full EP run, and it reaches an FID of 3.51 against the full architecture’s 3.54: the two are indistinguishable at the resolution of the measurement. Generation competitive with the diffusion and flow-matching baselines of Table 2 therefore does not require either component. This architecture is wider (78M against 50M) and used a different schedule and a longer first phase (App. I), so we take it as an existence proof and leave a cheaper recipe to future work.

6 RELATED WORK

Predictive coding. Generative predictive coding typically places the image at the bottom of a hierarchy of predictions: Monte Carlo predictive coding samples it by adding Langevin noise to inference, at MNIST scale (Oliviers et al., 2024), and Langevin predictive coding samples a latent code the same way and decodes it, reaching CIFAR-10 at FID 113 (Zahid et al., 2023). In both, the image is the predicted layer of a prediction-error likelihood. In ours it is the input of a network whose scalar output is an explicit potential, and the force on it is that potential’s gradient.

Generation by energy descent. Energy Matching (Balcerak et al., 2025), which we build on, generates by descent on an explicit scalar potential. Equilibrium Matching (Wang & Du, 2025) learns a time-invariant gradient field, in its best-performing form without an explicit energy, and Rodriguez et al. (2026) train it on an Energy Transformer, whose blocks relax latent tokens on a Hopfield-type energy, reaching class-conditional FID 28.56 on CIFAR-10. There the image-space update is a learned function of the relaxed tokens rather than the gradient of any energy in x ; in ours it is the gradient of $\gamma^{-1}E_{\text{PCN}} = V_{\theta} + O(\gamma)$. Each of these works trains by backpropagation.

Physical and neuromorphic generation. A parallel line generates with physical substrates. Simulated Kuramoto oscillators (Unconventional AI, 2026) and continuous spins descending a learned quadratic energy (Sun et al., 2026) evolve a latent state from a random start and render it through a small convolutional decoder, reaching class-conditional CIFAR-10 FID 8.76 and clean-FID 9.71 respectively, the former also at ImageNet 64×64 . Optical systems have implemented a diffusion model’s denoiser (Oguz et al., 2024) and a single-pass generator (Chen et al., 2025), demonstrated in hardware at MNIST scale, and spiking diffusion models target event-driven neuromorphic chips (Cao et al., 2024). All of these are trained by backpropagation. Bösch et al. (2026) generate 28×28 MNIST zeros and ones with a local rule, on an energy with no hidden units. Ours relaxes in image space with no decoder, and uses Equilibrium Propagation to obtain local updates for an energy with hidden layers.

7 DISCUSSION AND LIMITATIONS

Simulation cost. On a GPU, an EP training step costs about sixteen times a backpropagation step, and an implicit-solve step about five, with the relaxation budgets dominating. A substrate that settles physically would not pay that cost, but its speed, energy and settling accuracy remain unmeasured: everything here is simulated.

Locality and hardware realisation. The relaxation is local in the sense EP requires: every quantity a weight update needs is present at its two ends. The nodes of our training runs are, however, architecture blocks rather than single weight matrices, chosen to reduce simulation cost, so a literal implementation would need circuitry to carry the error through each block’s internal layers. One node per weight matrix removes this, and left the training gradient unchanged in the one checkpoint we measured (App. D). Analog EP also obtains the backward drive $J^{\top}e$ from a feedback pathway programmed to match the forward one; any drift between the two leaves dynamics that descend no energy at all (Kendall et al., 2020).

Scope of the evidence. App. K collects the scope of each claim. Each arm is a single run, so training variance is unknown, and the small FID differences do not establish a ranking. We demonstrate

EP training on CIFAR-10 only; at ImageNet-32 we verify the inference correspondence on the authors' released checkpoints, but did not train. The hardware-plausible architecture was trained and sampled only as a feedforward network, so it has not been demonstrated together with EP. The equilibria were computed by noise-free digital relaxation in error coordinates; the local state dynamics share their fixed points (App. E) but were not run to equilibrium at this scale. Analog noise, finite bit depth and forward-backward pathway mismatch remain untested.

AI USE STATEMENT

We used a frontier large language model in three ways. For the paper’s prose, we used it iteratively at the sentence level — for wording, concision and clarity — rather than to draft sections or generate content; the ideas, claims and structure are the authors’. For the code, we used it to help write and debug the training, sampling and measurement implementations; every suggestion was read and accepted or rejected manually, and the numerical results reported here come from runs of the accepted code. For related work, we used it to expand a literature search the authors had already conducted; we checked every claim we attribute to a cited work against that work directly. We did not use it to generate research ideas or hypotheses, to analyse results, or to produce figures. We have reviewed all AI-assisted work and take responsibility for the final content of this paper, including its text, claims and artifacts.

REPRODUCIBILITY STATEMENT

The PCN construction and the correspondence it rests on are given in Section 3.1, with the full derivation, its assumptions and the order of every neglected term in App. A; the extension to the skip-connected UNet is in App. D and the solver we relax with in App. E. The training construction, its parameters and the contrastive phase are in Section 3.2 and App. B. App. J gives every training and sampling setting, the FID protocol, the measurement resolution, and the implementation choices the PCN formulation requires but a feedforward one does not; App. I specifies the norm- and attention-free architecture and its recipe. Anonymised code for training, sampling and the correspondence measurements is available at <https://anonymous.4open.science/r/PotentialEP>.

REFERENCES

- Michal Balcerak, Tamaz Amiranashvili, Antonio Terpin, Suprosanna Shit, Lea Bogensperger, Sebastian Kaltenbach, Petros Koumoutsakos, and Bjoern Menze. Energy Matching: Unifying Flow Matching and Energy-Based Models for Generative Modeling, October 2025. URL <http://arxiv.org/abs/2504.10612>. arXiv:2504.10612 [cs].
- Cyrill Bösch, Geoffrey Roeder, Marc Serra-Garcia, and Ryan P. Adams. Local Learning Rules for Out-of-Equilibrium Physical Generative Models, July 2026. URL <http://arxiv.org/abs/2506.19136>. arXiv:2506.19136 [cs.LG].
- Jiahang Cao, Ziqing Wang, Hanzhong Guo, Hao Cheng, Qiang Zhang, and Renjing Xu. Spiking Denoising Diffusion Probabilistic Models. In *2024 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pp. 4900–4909, Waikoloa, HI, USA, January 2024. IEEE. ISBN 979-8-3503-1892-0. doi: 10.1109/WACV57701.2024.00484. URL <https://ieeexplore.ieee.org/document/10484223/>.
- Shiqi Chen, Yuhang Li, Hanlong Chen, and Aydogan Ozcan. Optical Generative Models. *Nature*, 644(8078):903–911, August 2025. ISSN 0028-0836, 1476-4687. doi: 10.1038/s41586-025-09446-5. URL <http://arxiv.org/abs/2410.17970>. arXiv:2410.17970 [cs.NE].
- Soham De and Samuel L. Smith. Batch Normalization Biases Residual Blocks Towards the Identity Function in Deep Networks, December 2020. URL <http://arxiv.org/abs/2002.10444>. arXiv:2002.10444 [cs.LG].
- Yilun Du, Shuang Li, Joshua Tenenbaum, and Igor Mordatch. Improved contrastive divergence training of energy-based models. In Marina Meila and Tong Zhang (eds.), *Proceedings of the 38th international conference on machine learning*, volume 139 of *Proceedings of machine learning research*, pp. 2837–2848. PMLR, July 2021. URL <https://proceedings.mlr.press/v139/du21b.html>.
- Cédric Goemaere, Gaspard Oliviers, Rafal Bogacz, and Thomas Demeester. ePC: Fast and Deep Predictive Coding in Digital Simulation, June 2026. URL <http://arxiv.org/abs/2505.20137>. arXiv:2505.20137 [cs.LG].

- Alex Gower. FlowEqProp: Training Flow Matching Models with Gradient Equilibrium Propagation. In *Proceedings of the International Conference on Neuromorphic Systems, ICONS '26*, pp. 60–66, New York, NY, USA, August 2026. Association for Computing Machinery. ISBN 979-8-4007-2808-2. doi: 10.1145/3822454.3822485. URL <https://doi.org/10.1145/3822454.3822485>.
- Rasmus Høier, D. Staudt, and Christopher Zach. Dual Propagation: Accelerating Contrastive Hebbian Learning with Dyadic Neurons. In *Proceedings of the 40th International Conference on Machine Learning*, pp. 13141–13156. PMLR, July 2023. URL <https://proceedings.mlr.press/v202/hoier23a.html>.
- Daniele Ielmini and H.-S. Philip Wong. In-memory computing with resistive switching devices. *Nature Electronics*, 1(6):333–343, June 2018. ISSN 2520-1131. doi: 10.1038/s41928-018-0092-2. URL <https://www.nature.com/articles/s41928-018-0092-2>.
- Jack Kendall, Ross Pantone, Kalpana Manickavasagam, Yoshua Bengio, and Benjamin Scellier. Training End-to-End Analog Neural Networks with Equilibrium Propagation, June 2020. URL <http://arxiv.org/abs/2006.01981>. arXiv:2006.01981 [cs].
- Tugdual Kerjan, Rasmus Høier, and Benjamin Scellier. Training a Predictive Coding Network on ImageNet using Equilibrium Propagation, June 2026. URL <http://arxiv.org/abs/2606.03584>. arXiv:2606.03584 [cs.LG].
- Dongjun Kim, Seungjae Shin, Kyungwoo Song, Wanmo Kang, and Il-Chul Moon. Soft Truncation: A Universal Training Technique of Score-based Diffusion Model for High Precision Score Estimation. In *Proceedings of the 39th International Conference on Machine Learning*, pp. 11201–11228. PMLR, June 2022. URL <https://proceedings.mlr.press/v162/kim22i.html>.
- Axel Laborieux, Maxence Ernoult, Benjamin Scellier, Yoshua Bengio, Julie Grollier, and Damien Querlioz. Scaling Equilibrium Propagation to Deep ConvNets by Drastically Reducing its Gradient Estimator Bias, June 2020. URL <http://arxiv.org/abs/2006.03824>. arXiv:2006.03824.
- Hankook Lee, Jongheon Jeong, Sejun Park, and Jinwoo Shin. Guiding Energy-based Models via Contrastive Latent Variables. In *ICLR*, September 2022. URL <https://openreview.net/forum?id=CZmHHj9MgkP>.
- Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow Matching for Generative Modeling, February 2023. URL <http://arxiv.org/abs/2210.02747>. arXiv:2210.02747 [cs].
- Sasha Luccioni, Yacine Jernite, and Emma Strubell. Power Hungry Processing: Watts Driving the Cost of AI Deployment? In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency, FAccT '24*, pp. 85–99, New York, NY, USA, June 2024. Association for Computing Machinery. ISBN 979-8-4007-0450-5. doi: 10.1145/3630106.3658542. URL <https://dl.acm.org/doi/10.1145/3630106.3658542>.
- Tom Van Der Meersch, Johannes Deleu, and Thomas Demeester. Training a Hopfield Variational Autoencoder with Equilibrium Propagation, November 2023. URL <http://arxiv.org/abs/2311.15047>. arXiv:2311.15047 [cs].
- Beren Millidge, Alexander Tschantz, and Christopher L. Buckley. Predictive Coding Approximates Backprop Along Arbitrary Computation Graphs. *Neural Computation*, 34(6):1329–1368, May 2022. ISSN 0899-7667. doi: 10.1162/neco.a.01497. URL <https://doi.org/10.1162/neco.a.01497>.
- Ilker Oguz, Niyazi Ulas Dinc, Mustafa Yildirim, Junjie Ke, Innfarn Yoo, Qifei Wang, Feng Yang, Christophe Moser, and Demetri Psaltis. Optical diffusion models for image generation. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (eds.), *Advances in neural information processing systems*, volume 37, pp. 59150–59173. Curran Associates, Inc., 2024. doi: 10.52202/079017-1887. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/6cb81234ab47027e991728ed7dd76735-Paper-Conference.pdf.

- Gaspard Oliviers, Rafal Bogacz, and Alexander Meulemans. Learning probability distributions of sensory inputs with Monte Carlo predictive coding. *PLOS Computational Biology*, 20(10):e1012532, October 2024. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1012532. URL <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1012532>.
- Siyuan Qiao, Huiyu Wang, Chenxi Liu, Wei Shen, and Alan Yuille. Micro-Batch Training with Batch-Channel Normalization and Weight Standardization, August 2020. URL <http://arxiv.org/abs/1903.10520>. arXiv:1903.10520 [cs.CV].
- Rajesh P. N. Rao and Dana H. Ballard. Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1):79–87, January 1999. ISSN 1546-1726. doi: 10.1038/4580. URL <https://doi.org/10.1038/4580>.
- Adrian Rodriguez, Benjamin Hoover, Yunhui Guo, and Yilun Du. Generative Associative Memory via Equilibrium Matching. In *ICLR*, March 2026. URL <https://openreview.net/forum?id=4zjhUaeeVh>.
- Benjamin Scellier and Yoshua Bengio. Equilibrium Propagation: Bridging the Gap Between Energy-Based Models and Backpropagation, March 2017. URL <http://arxiv.org/abs/1602.05179>. arXiv:1602.05179.
- Abu Sebastian, Manuel Le Gallo, Riduan Khaddam-Aljameh, and Evangelos Eleftheriou. Memory devices and applications for in-memory computing. *Nature Nanotechnology*, 15(7):529–544, July 2020. ISSN 1748-3395. doi: 10.1038/s41565-020-0655-z. URL <https://www.nature.com/articles/s41565-020-0655-z>.
- Emma Strubell, Ananya Ganesh, and Andrew McCallum. Energy and Policy Considerations for Deep Learning in NLP. In Anna Korhonen, David Traum, and Lluís Màrquez (eds.), *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 3645–3650, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1355. URL <https://aclanthology.org/P19-1355/>.
- Lianlong Sun, Chuan Liu, Tong Geng, and Michael Huang. Conditional Dynamical Systems for Image Generation, August 2026. URL <http://arxiv.org/abs/2608.14961>. arXiv:2608.14961 [cs.ET].
- Alexander Tong, Kilian Fatras, Nikolay Malkin, Guillaume Hugué, Yanlei Zhang, Jarrod Rector-Brooks, Guy Wolf, and Yoshua Bengio. Improving and generalizing flow-based generative models with minibatch optimal transport, March 2024. URL <http://arxiv.org/abs/2302.00482>. arXiv:2302.00482 [cs].
- Unconventional AI. Introducing Un-0: Generating Images with Coupled Oscillators - Unconventional AI, 2026. URL <https://unconv.ai/blog/introducing-un-0-generating-images-with-coupled-oscillators/>.
- Runqian Wang and Yilun Du. Equilibrium Matching: Generative Modeling with Implicit Energy-Based Models, October 2025. URL <http://arxiv.org/abs/2510.02300>. arXiv:2510.02300 [cs.LG].
- James C. R. Whittington and Rafal Bogacz. An Approximation of the Error Backpropagation Algorithm in a Predictive Coding Network with Local Hebbian Synaptic Plasticity. *Neural Computation*, 29(5):1229–1262, May 2017. ISSN 1530-888X. doi: 10.1162/NECO_a.00949.
- Logan G. Wright, Tatsuhiko Onodera, Martin M. Stein, Tianyu Wang, Darren T. Schachter, Zoey Hu, and Peter L. McMahon. Deep physical neural networks trained with backpropagation. *Nature*, 601(7894):549–555, January 2022. ISSN 1476-4687. doi: 10.1038/s41586-021-04223-6. URL <https://www.nature.com/articles/s41586-021-04223-6>.
- Shu Yang, Hanzhi Ma, Chengting Yu, Aili Wang, Pengqi Fu, and Er-Ping Li. SDiT: Scalable and efficient spiking diffusion transformers for image generation. *Pattern Recogn. Lett.*, 207(C): 266–273, August 2026. ISSN 0167-8655. doi: 10.1016/j.patrec.2026.07.003. URL <https://doi.org/10.1016/j.patrec.2026.07.003>.

Umais Zahid, Qinghai Guo, and Zafeirios Fountas. Sample as You Infer: Predictive Coding With Langevin Dynamics, November 2023. URL <http://arxiv.org/abs/2311.13664>. arXiv:2311.13664 [cs.LG] version: 1.

A DERIVATION OF THE CORRESPONDENCE

This appendix fills in Proposition 1: the assumptions, the order of every neglected term, and where the $O(\gamma)$ error comes from. Throughout, α is a fixed constant of the architecture and is absorbed into the $O(\cdot)$ symbols, so “ $O(\gamma)$ ” means $O(\alpha\gamma)$.

A.1 THE ARGUMENT IN OUTLINE

Note first that without the injection the energy is minimised at $e = 0$, that is, at the forward pass $h_k^{\text{FFN}} = (f_k \circ \dots \circ f_1)(x)$ for every k . The injection disturbs this. Stationarity in the output node gives $\partial E / \partial h_N = -e_N + \gamma\alpha = 0$, so the top error is a constant, $e_N^* = \gamma\alpha$, exactly. Stationarity in an interior node gives $\partial E / \partial h_k = -e_k + J_{k+1}^\top e_{k+1} = 0$ with $J_{k+1} = \partial f_{k+1} / \partial h_k$, so

$$e_k^* = J_{k+1}^\top e_{k+1}^*, \quad \text{hence} \quad e_1^* = \gamma\alpha J_2^\top J_3^\top \dots J_N^\top, \quad (13)$$

with every Jacobian evaluated at h^* . Each error is therefore the derivative of the output with respect to its own node, propagated down by transposed Jacobians from a seed of $\gamma\alpha$ at the top — the same recursion backpropagation performs, seeded there by $\partial V / \partial V = 1$. Every error is $O(\gamma)$, so from $h_k^* = f_k(h_{k-1}^*) - e_k^*$ each node h_k^* sits $O(\gamma)$ from the forward pass state h_k^{FFN} it would occupy at $\gamma = 0$.

The velocity needs the total derivative of $E_{\text{PCN}}(x, h^*(x))$ with respect to x . Only the first error term depends on x explicitly, contributing $\partial E / \partial x = J_1^\top e_1^*$; the dependence through $h^*(x)$ adds $\sum_k (\partial E / \partial h_k) \big|_{h^*} (dh_k^* / dx)$, which vanishes term by term because h^* is stationary. Hence

$$\nabla_x E_{\text{PCN}}(x, h^*(x)) = J_1^\top e_1^* = \gamma\alpha J_1^\top J_2^\top \dots J_N^\top = \gamma\alpha \nabla_x F_\theta(x) + O(\gamma^2), \quad (14)$$

the last step evaluating each Jacobian at h^{FFN} rather than at h^* , which costs $O(\gamma)$ each, and recognising the resulting product as the chain rule for $\nabla_x F_\theta$. Dividing by γ and using $V_\theta = \alpha F_\theta$ gives the first claim. For the second, evaluate E_{PCN} at h^* : the error terms are $O(\gamma^2)$ by equation 13, and $o^* = f_N(h_{N-1}^*) - \gamma\alpha = F_\theta(x) + O(\gamma)$, so the injection term contributes $\gamma V_\theta(x) + O(\gamma^2)$.

A.2 THE ARGUMENT IN DETAIL

This section repeats the steps above with the assumptions stated and every neglected term accounted for.

Setting. E_{PCN} is equation 6 on a chain of N nodes with $h_0 = x$ and $h_N = o$ scalar. Each f_k is C^2 in its input and in its parameters θ , with bounded first and second derivatives on the relevant region (SiLU, tanh; ReLU almost everywhere). The differentiability in θ that Section 3.2 requires needs C^4 , which smooth activations have. Write $h_k^{\text{FFN}} = (f_k \circ \dots \circ f_1)(x)$ for the forward-pass state at node k , with $h_0^{\text{FFN}} = x$ and $h_N^{\text{FFN}} = F_\theta(x)$, and $J_k(h) = \partial f_k / \partial h$ for the Jacobian of layer k at its input. The chain rule gives $\nabla_x F_\theta(x) = J_1^\top J_2^\top \dots J_N^\top$, where each J_k is evaluated at the input layer k receives on the forward pass, namely h_{k-1}^{FFN} ; note J_N is a row vector, so the product is a column of dimension $\dim x$.

Existence of the equilibrium. At $\gamma = 0$, $E_{\text{PCN}} \geq 0$ with a unique minimum at $e = 0$, i.e. $h_k = f_k(h_{k-1})$ for every k , which by induction from $h_0 = x$ is the forward pass $h_k = h_k^{\text{FFN}}$. In the error coordinates of App. E the Hessian there is the identity, so by the implicit function theorem the minimiser continues smoothly to small $\gamma > 0$, and depends smoothly on θ and γ jointly; $h^*(x)$ denotes that branch. Physically it is the state the hidden layers settle to when they relax faster than x moves.

Step 1: the top error. h_N enters E_{PCN} through $\frac{1}{2}(f_N(h_{N-1}) - h_N)^2$ and $\gamma\alpha h_N$, so $\partial E/\partial h_N = -e_N + \gamma\alpha$, and stationarity gives

$$e_N^* = \gamma\alpha \quad \text{exactly.} \quad (15)$$

This is a constant independent of x and θ : the injection seeds the backward pass with a fixed scalar.

Step 2: interior errors backpropagate. For $1 \leq k < N$, h_k enters its own error term and the next node's, giving $\partial E/\partial h_k = -e_k + J_{k+1}(h_k)^\top e_{k+1}^*$. Stationarity gives $e_k^* = J_{k+1}(h_k^*)^\top e_{k+1}^*$, and iterating from the top,

$$e_k^* = \gamma\alpha J_{k+1}(h_k^*)^\top \cdots J_N(h_{N-1}^*)^\top, \quad k = 1, \dots, N-1, \quad (16)$$

exact, with every Jacobian at the equilibrium state. Bounded Jacobians make every $e_k^* = O(\gamma)$.

Step 3: states sit $O(\gamma)$ from the forward pass. From $h_k^* = f_k(h_{k-1}^*) - e_k^*$ and $h_0 = x$, each node's displacement from the forward pass is its own $O(\gamma)$ error plus whatever its predecessor's displacement becomes after passing through f_k . Bounded first derivatives make f_k Lipschitz with constant $L_k = \sup \|J_k\|$, so the displacement compounds as $\|h_k^* - h_k^{\text{FFN}}\| \leq L_k \|h_{k-1}^* - h_{k-1}^{\text{FFN}}\| + \|e_k^*\|$ rather than growing without control, so $h_k^* = h_k^{\text{FFN}} + O(\gamma)$ at every depth. Bounded second derivatives then carry that displacement into the Jacobians. Since J_{k+1} is a function of the point at which it is evaluated, Taylor expanding about h_k^{FFN} gives

$$J_{k+1}(h_k^*) = J_{k+1}(h_k^{\text{FFN}}) + \nabla^2 f_{k+1}(h_k^{\text{FFN}})(h_k^* - h_k^{\text{FFN}}) + \dots \quad (17)$$

and the displacement in the second term is $O(\gamma)$ by the above, so $J_{k+1}(h_k^*) = J_{k+1}(h_k^{\text{FFN}}) + O(\gamma)$.

Step 4: the visible gradient. $x = h_0$ appears in E_{PCN} only through $\frac{1}{2}\|f_1(x) - h_1\|^2$, so at fixed h , $\partial_x E_{\text{PCN}} = J_1(x)^\top e_1$. The total derivative of $E_{\text{PCN}}(x, h^*(x))$ is $\partial_x E|_{h^*} + \sum_k (\partial_{h_k} E)|_{h^*} (dh_k^*/dx)$, and the second sum vanishes because h^* is a stationary point of E_{PCN} in h . Hence

$$\nabla_x E_{\text{PCN}}(x, h^*(x)) = J_1(x)^\top e_1^* = \gamma\alpha J_1(x)^\top J_2(h_1^*)^\top \cdots J_N(h_{N-1}^*)^\top. \quad (18)$$

Replacing each $J_k(h_{k-1}^*)$ by $J_k(h_{k-1}^{\text{FFN}}) + O(\gamma)$ and collecting, the product is $\nabla_x F_\theta(x) + O(\gamma)$, so

$$\nabla_x E_{\text{PCN}}(x, h^*(x)) = \gamma\alpha \nabla_x F_\theta(x) + O(\gamma^2) = \gamma \nabla_x V_\theta(x) + O(\gamma^2), \quad (19)$$

and $v_{\text{PCN}} = -\gamma^{-1} \nabla_x E_{\text{PCN}}(x, h^*) = -\nabla_x V_\theta(x) + O(\gamma)$.

Step 5: the energy value. At equilibrium the error terms contribute $\sum_k \frac{1}{2} \|e_k^*\|^2 = O(\gamma^2)$. For the injection term, Step 1 gives $e_N^* = \gamma\alpha$ and Step 3 gives $h_{N-1}^* = h_{N-1}^{\text{FFN}} + O(\gamma)$, so $o^* = f_N(h_{N-1}^*) - e_N^* = F_\theta(x) + O(\gamma)$ and $E_{\text{PCN}}(x, h^*) = \gamma\alpha F_\theta(x) + O(\gamma^2) = \gamma V_\theta(x) + O(\gamma^2)$.

Where the error lives. The stationarity conditions of Steps 1 and 2 and the envelope identity of Step 4 are exact, and J_1 is evaluated at x itself, which is held fixed and carries no displacement. The only approximation is Step 3's replacement of the interior Jacobians $J_2, \dots, J_N$ by their values on the forward pass, so the $O(\gamma)$ correction to v_{PCN} comes entirely from the interior states being displaced by their own errors: it vanishes for linear layers and is bounded by their second derivatives. The output node's own displacement plays no part — it has no children, and e_N^* enters everything downstream only through the exact constant $\gamma\alpha$. Section 4.1 measures the correction as linear in γ over six decades, as this predicts.

Directed acyclic graphs. On the skip-connected UNet of App. D, a node may have several parents and several children. Step 2 becomes $e_j^* = \sum_{c \in \text{ch}(j)} J_{c,j}^\top e_c^*$ with $J_{c,j} = \partial f_c / \partial h_j$, which is exactly the chain rule's multivariable form, and Step 4 is unchanged: every node fed by x contributes $J_{c,0}^\top e_c^*$ to the visible gradient. The rest of the argument goes through verbatim.

Differentiability in θ . Proposition 1 bounds the gap between v_{PCN} and $-\nabla_x V_\theta$ at fixed θ . The gradient correspondence of Section 3.2, on which training rests, needs more: that the gap’s θ -derivative is $O(\gamma)$ too, which does not follow from a pointwise bound alone. It follows instead from smoothness, in three steps.

First, the equilibrium is as smooth as the layers. Applying the implicit function theorem to the spring-clamped system equation 9 in the joint variables (x, e) , whose Hessian at $\gamma = 0$ is $\lambda I \oplus I$ and hence invertible, the free equilibrium (x^*, h^*) is a C^k function of (θ, γ) whenever the layers are C^{k+1} .

Second, the velocity is finite at $\gamma = 0$. Every error vanishes there, so $\nabla_x E_{\text{PCN}}(x^*, h^*)$ vanishes and $x^* = x_t$; a C^k function that vanishes at $\gamma = 0$, divided by γ , is C^{k-1} , so $v^* = -\gamma^{-1} \nabla_x E_{\text{PCN}}$ is C^{k-1} in (θ, γ) including at $\gamma = 0$, where by Step 4 its value is $-\nabla_x V_\theta(x_t)$. Hence $\mathcal{L}_{\text{PCN}}(\theta, 0) = \mathcal{L}(\theta)$.

Third, the mean value theorem in γ : if $\nabla_\theta \mathcal{L}_{\text{PCN}}$ has a bounded γ -derivative, its change from $\gamma = 0$ is at most that bound times γ . That derivative is a mixed second derivative of $\mathcal{L}_{\text{PCN}}$, so $\mathcal{L}_{\text{PCN}}$ — and hence v^* , of which it is a smooth function — must be jointly C^2 . Tracing back through the two derivatives lost above, one to the implicit function theorem and one to the division by γ , this holds for layers with four continuous derivatives, which smooth activations such as SiLU and tanh have. Then $\nabla_\theta \mathcal{L}_{\text{PCN}}(\theta, \gamma) = \nabla_\theta \mathcal{L}(\theta) + O(\gamma)$, which is the correspondence.

B THE TRAINING CONSTRUCTION IN DETAIL

This appendix fills in the construction of Section 3.2.

$\mathcal{L}_{\text{PCN}}$ differs from $\mathcal{L}$ in two ways, both $O(\gamma)$. First, the correspondence is inexact at finite γ : $v_{\text{PCN}} = -\nabla_x V_\theta + O(\gamma)$ by Proposition 1. Second, $v^* = v_{\text{PCN}}(x^*)$ is the velocity at the relaxed state x^* rather than at x_t , an evaluation point displaced by $x^* - x_t = \gamma v^* / \lambda$. Keeping that displacement small is the tether λ ’s job, and it has less work to do here than in FlowEqProp (Gower, 2026), where GradEP was originally introduced: there the velocity converted to an extension with a factor $O(1/\lambda)$, so a close evaluation point required $\lambda \gg 1$, whereas equation 10 converts with γ/λ so the requirement is only $\gamma/\lambda \ll 1$. The smallness of γ that the correspondence already demands therefore supplies this constraint too, leaving λ free to be chosen on other grounds (see below).

γ and β . The injection strength γ plays for V the role β plays for the cost C in Equilibrium Propagation: a small linear term on one node whose stationary response, divided by that term’s strength, is a gradient. As in EP, the division is a constant gain on the readout, which costs nothing on an analog substrate.

Sweep budgets. The free phase relaxes for T_{free} sweeps and each nudged phase for T_{nudge} sweeps from the free equilibrium, with h equilibrated by K_h inner sweeps per step of x in both.

Parameters. Five quantities appear in the training construction besides γ . The spring stiffness λ is not tuned: with the visible state updated as $x \leftarrow x - dt_x \nabla_x E_{\text{free}}$, the spring extension contracts by $(1 - dt_x \lambda)$ per sweep, up to the $O(\gamma)$ curvature of E_{PCN} in x (App. C), so $\lambda dt_x = 1$ settles it in a single step, exactly as the error parameterisation settles h in one sweep; larger λdt_x overshoots. Phase 1 uses $\lambda = 1$, $dt_x = 1$ and phase 2 uses $\lambda = 2$, $dt_x = 0.5$ (App. J). The nudge strength β is the one parameter that must be chosen: it sets the displacement between the two nudged equilibria, which must be large enough to resolve in finite-precision arithmetic and small enough for equation 11 to hold, a window Section 3.2.1 derives and Section 4.2 measures. The budgets T_{free} , T_{nudge} and K_h are the sweep counts of the three relaxations; the last two are satisfied by a single sweep for the same deadbeat reason as λ , so the cost of an EP update is the free-phase equilibration T_{free} alone.

Contrastive divergence. Energy Matching’s second phase adds to the flow-matching loss a contrastive term, $\nabla_\theta V(x^+) - \nabla_\theta V(x^-)$ between data and model samples. On the PCN this needs no nudge at all. Writing $V = \gamma^{-1} E_{\text{PCN}}(x, h^*(x; \theta))$, the parameter derivative of E_{PCN} is $\partial_\theta E_{\text{PCN}}|_{h^*} + \sum_k \partial_{h_k} E_{\text{PCN}}|_{h^*} (dh_k^*/d\theta)$, whose second term vanishes at equilibrium. What re-

mains is local: for a parameter in layer k , $\partial E_{\text{PCN}}/\partial \theta_k = e_k^\top \partial f_k/\partial \theta_k$, involving only the error at node k and the state feeding it. The contrastive gradient is therefore the difference of $\gamma^{-1} \partial_\theta E_{\text{PCN}}|_{h^*}$ between the hidden equilibria reached with x clamped at x^+ and at x^- — a contrastive Hebbian rule.

C FINITE-DIFFERENCE ACCURACY OF THE LINEARISED NUDGE

Write $d = x^* - x_t$ and $d_\beta = x^{*\beta} - x_t$ for the free and nudged extensions, so that $\nabla_x E_{\text{PCN}}(x^*) = -\lambda d$ is the force balance at the free equilibrium. Moving x changes $\nabla_x E_{\text{PCN}}$ by $\nabla_x^2 E_{\text{PCN}}$ times the displacement, and $\nabla_x^2 E_{\text{PCN}} = O(\gamma)$: the gradient $\nabla_x E_{\text{PCN}} = J_1^\top e_1$ is $O(\gamma)$ in size (App. A, Step 2), and the distance in x over which it changes appreciably is set by the $O(1)$ weights and nonlinearities rather than by γ , so each of its x -derivatives is $O(\gamma)$ too. The tether’s Hessian is $\nabla_x^2 \frac{\lambda}{2} \|x - x_t\|^2 = \lambda I$, so across the same displacement $\nabla_x E_{\text{PCN}}$ varies by a fraction $O(\gamma/\lambda)$ of what the tether does: we treat $\nabla_x E_{\text{PCN}}$ as constant at its free-phase value $-\lambda d$ throughout the nudge, while keeping the full variation of the tether and nudge terms.

Stationarity of $E_{\text{free}} + \beta' C$ in x then sets three gradient terms against each other, and the nudge term differs between the two costs.

Quadratic cost. With $v(x) = \frac{\lambda}{\gamma}(x - x_t)$,

$$\beta' C(x) = \beta \left(\frac{\gamma}{\lambda} \right)^2 \frac{1}{2} \left\| \frac{\lambda}{\gamma}(x - x_t) - u \right\|^2, \quad \nabla_x \beta' C = \beta \left(x - x_t - \frac{\gamma}{\lambda} u \right) = \beta \left(d_\beta - \frac{\gamma}{\lambda} u \right), \quad (20)$$

so stationarity reads

$$\underbrace{-\lambda d}_{\nabla_x E_{\text{PCN}}} + \underbrace{\lambda d_\beta}_{\text{tether}} + \underbrace{\beta \left(d_\beta - \frac{\gamma}{\lambda} u \right)}_{\text{nudge}} = 0 \implies d_\beta = \frac{\lambda d + \beta \gamma u / \lambda}{\lambda + \beta}, \quad (21)$$

and subtracting d gives

$$d_\beta - d = \frac{\beta(\gamma u / \lambda - d)}{\lambda + \beta}, \quad (22)$$

into which $d = \gamma v^* / \lambda$ from equation 10 and $g = v^* - u$ substitute to give

$$x^{*\beta} - x^* = d_\beta - d = -\frac{\beta \gamma g}{\lambda(\lambda + \beta)}. \quad (23)$$

This is nonlinear in β through the $\lambda + \beta$ in the denominator, and linear only while $\beta \ll \lambda$. Replacing β by $-\beta$ gives a displacement of different magnitude, $\lambda - \beta$ in place of $\lambda + \beta$, so the $\pm\beta$ equilibria are not symmetric about x^* , and the centred estimator equation 11 inherits a bias the next paragraphs put at $O((\beta/\lambda)^2)$.

Linearised cost. This targets the same gradient as the quadratic cost: g is a constant during the nudged phases, so $\nabla_\theta C_{\text{lin}}(x^*) = g^\top \nabla_\theta v^*$, which is the chain rule for $\nabla_\theta \frac{1}{2} \|v^* - u\|^2$.

Here

$$\beta' C_{\text{lin}}(x) = \beta \left(\frac{\gamma}{\lambda} \right)^2 \frac{\lambda}{\gamma} g^\top (x - x_t) = \beta \frac{\gamma}{\lambda} g^\top (x - x_t), \quad \nabla_x \beta' C_{\text{lin}} = \beta \frac{\gamma}{\lambda} g, \quad (24)$$

a constant force independent of x , so stationarity reads

$$\underbrace{-\lambda d}_{\nabla_x E_{\text{PCN}}} + \underbrace{\lambda d_\beta}_{\text{tether}} + \underbrace{\beta \frac{\gamma}{\lambda} g}_{\text{nudge}} = 0 \implies d_\beta = d - \frac{\beta \gamma g}{\lambda^2}, \quad (25)$$

and subtracting d gives

$$x^{*\beta} - x^* = d_\beta - d = -\frac{\beta \gamma g}{\lambda^2}. \quad (26)$$

There is no β in the denominator: the displacement is linear in β , so the $\pm\beta$ equilibria are symmetric about x^* up to the $O(\gamma/\lambda)$ accuracy of holding $\nabla_x E_{\text{PCN}}$ fixed. The other β -dependent error that remains comes from $\partial_\theta E_{\text{PCN}}$ varying nonlinearly along the straight path between the two equilibria, which the next paragraphs bound.

The centred estimate. Write $\Phi(x) = \partial_\theta E_{\text{PCN}}(x, h^*(x))$ for the local gradient at the hidden equilibrium for a given visible state, and $G(\beta) = \Phi(x^{*\beta})$. Expanding $G(\pm\beta)$ about $\beta = 0$ and subtracting cancels the even orders,

$$\frac{G(\beta) - G(-\beta)}{2\beta} = G'(0) + \frac{\beta^2}{6} G'''(0) + O(\beta^4), \quad (27)$$

so the estimate in equation 11 returns the wanted $G'(0)$ with a relative error $\beta^2 |G'''(0)| / 6 |G'(0)|$, componentwise. What differs between the two costs is where that third derivative comes from.

Nonlinearity in β under the quadratic cost. From equation 23 the nudged state moves along the fixed direction $-g$ at a rate nonlinear in β ,

$$x^{*\beta} = x^* + f(\beta) w_0, \quad w_0 = -\frac{\gamma g}{\lambda}, \quad f(\beta) = \frac{\beta}{\lambda + \beta}, \quad (28)$$

and by the chain rule

$$G'(0) = f'(0) \nabla_x \Phi \cdot w_0, \quad (29)$$

$$G'''(0) = f'''(0) \nabla_x \Phi \cdot w_0 + 3f'(0)f''(0) \nabla_x^2 \Phi[w_0, w_0] + f'(0)^3 \nabla_x^3 \Phi[w_0, w_0, w_0]. \quad (30)$$

Three things can bias the estimate: the nonlinearity of f , giving the first term of $G'''(0)$; the curvature of Φ , giving the other two, which carry its second and third derivatives; and the approximation of the previous section that held $\nabla_x E_{\text{PCN}}$ fixed, a relative $O(\gamma/\lambda)$ error in the equilibria themselves. The last does not grow with β and so cannot set a ceiling on it. Of the other two, the first term shares the factor $\nabla_x \Phi \cdot w_0$ with $G'(0)$, so

$$\frac{\beta^2}{6} \frac{f'''(0) \nabla_x \Phi \cdot w_0}{f'(0) \nabla_x \Phi \cdot w_0} = \frac{\beta^2}{6} \frac{6/\lambda^3}{1/\lambda} = \left(\frac{\beta}{\lambda}\right)^2, \quad (31)$$

a relative error present for any Φ and at any γ . The two Φ -curvature terms of $G'''(0)$ are smaller than the first by $\|w_0\|$ and $\|w_0\|^2$ respectively: each extra w_0 costs a factor $\|w_0\| = \gamma|g|/\lambda$, while the $O(\gamma)$ amplitudes of the $\nabla_x^n \Phi$ themselves are common to all three terms and cancel (next paragraph). Those ratios are independent of β , and at the small γ the correspondence requires they are tiny — $\|w_0\| \approx 4 \times 10^{-5}$ at the settings of Section 4.2 — so under the quadratic cost it is f 's nonlinearity that limits β . This is the asymmetry of the $\pm\beta$ equilibria noted above, now quantified: requiring $(\beta/\lambda)^2 \ll 1$ is the constraint $\beta \ll \lambda$ of equation 12.

Nonlinearity in β under the linearised cost. From equation 26 the path is $x^{*\beta} = x^* + f(\beta) w_0$ with the same direction $w_0 = -\gamma g/\lambda$, but now $f(\beta) = \beta/\lambda$ exactly. Every derivative of f beyond the first vanishes, so the first two terms of $G'''(0)$ go with them and

$$G'(0) = \frac{1}{\lambda} \nabla_x \Phi \cdot w_0, \quad G'''(0) = \frac{1}{\lambda^3} \nabla_x^3 \Phi[w_0, w_0, w_0]. \quad (32)$$

Only Φ 's own curvature is left. Its size follows from the same argument as before: every error at the hidden equilibrium is $O(\gamma)$ (App. A, Step 2) and Φ is bilinear in errors and activations, so Φ has amplitude $O(\gamma)$, while the distance in x over which it changes appreciably is set by the $O(1)$ weights and nonlinearities, so $\|\nabla_x^n \Phi\| = O(\gamma)$ for every n . Those $O(\gamma)$ amplitudes cancel in the ratio, leaving two extra factors of w_0 and two of λ :

$$\frac{\beta^2 |G'''(0)|}{6 |G'(0)|} = O\left(\frac{\beta^2 \|w_0\|^2}{\lambda^2}\right) = O\left(\left(\frac{\beta \gamma |g|}{\lambda^2}\right)^2\right) = O(|x^{*\beta} - x^*|^2) : \quad (33)$$

the estimate is accurate while the nudge displacement is small on the scale over which Φ varies. Unlike the quadratic cost's bias this is a condition on the displacement, which smaller γ relaxes rather than leaves untouched, and requiring it to be small is the constraint $\beta \gamma |g| \ll \lambda^2$ of equation 12.

Consistency with measurement. Both displacement formulas reproduce the diagnostics logged during training to three significant figures. At the operating point ($\gamma = 10^{-4}$, $\lambda = 1$, $\beta = 30$), with the logged per-element RMS values $|v^*| = 1.03$ and $|g| = 0.39$, equation 10 predicts a free-phase extension $|x^* - x_t| = 1.03 \times 10^{-4}$ and equation 26 a nudge displacement of 1.16×10^{-3} ; the logged values are 1.03×10^{-4} and 1.16×10^{-3} . The two constraints then place the ceilings four orders of magnitude apart: $\beta \ll \lambda = 1$ for the quadratic cost, against $\beta \ll \lambda^2/(\gamma|g|) = 2.6 \times 10^4$ for the linearised one. Section 4.2 measures the quadratic nudge failing above $\beta = 1$ and the linear nudge degrading only gently towards 10^4 .

Scope. The linearisation does not depend on the form of E_{PCN} , so it applies to any use of GradEP with a quadratic loss. It is also simpler to realise physically: a constant force on the visible nodes, rather than an added quadratic potential whose restoring force must track the displacement.

D PREDICTIVE CODING ON A DAG

D.1 FROM CHAINS TO DAGS

Predictive coding networks are usually written on a *chain*, with node k predicted from node $k - 1$ alone:

$$E_{\text{chain}} = \sum_{k=1}^N \frac{1}{2} \|f_k(h_{k-1}) - h_k\|^2, \quad (34)$$

with h_0 the clamped input. The UNet we adopt is not a chain: its decoder blocks consume a concatenation of their predecessor and a topologically distant encoder state, so a node may have several parents, and an encoder node several children. Predictive coding has been formulated on arbitrary computation graphs (Millidge et al., 2022), and we use that construction unchanged: nodes $h_1, \dots, h_N$ are taken in topological order, each with a parent list $\text{pa}(k)$, and

$$E = \sum_{k=1}^N \frac{1}{2} \|f_k(h_{\text{pa}(k)}) - h_k\|^2 + \gamma \alpha o, \quad (35)$$

with $o = h_N$ the scalar output node and $\gamma \alpha$ the injection of Section 3.1. The prediction functions are the unmodified architecture blocks; where a node has several parents they are concatenated in the order of the feedforward pass, so f_k evaluates the same function of the same inputs.

D.2 RELAXATION DYNAMICS

Writing $e_k = f_k(h_{\text{pa}(k)}) - h_k$ and $J_{c,j} = \partial f_c / \partial h_j$, each state evolves under

$$\frac{\partial E}{\partial h_j} = -e_j + \sum_{c \in \text{ch}(j)} J_{c,j}^\top e_c + \gamma \alpha \delta_{jN}, \quad (36)$$

the node’s own error term, the error term of every child that consumes it, and the output injection. In a chain $\text{ch}(j) = \{j + 1\}$; in the DAG a node with fan-out is driven by all its children at once, and stationarity is the multivariable chain rule (App. A). Every term lives on an edge of the graph, so locality is preserved.

D.3 NODE GRANULARITY

How a network is partitioned into nodes is a free choice. A node can be an architecture *block* — one residual block, one attention block, or the ViT head etc — in which case only the block’s output relaxes and its internal activations are computed feedforward inside f_k ; or it can be a single weight matrix and its activation, in which case every hidden activation relaxes. The correspondence of Section 3.1 holds at any partition. What the choice sets is how the backward pass is divided between relaxation and circuitry: finer partitions relax more states and need less backward circuitry, which suits hardware; coarser ones relax fewer states and need more backpropagation, which suits simulation. For a parameter inside node k ,

$$\frac{\partial E}{\partial \theta} = e_k^\top \frac{\partial f_k}{\partial \theta}. \quad (37)$$

At one node per weight matrix, $f_c(h_j) = \sigma(W h_j)$, the update is an outer product of quantities at the two ends of the weight, $\partial E / \partial W = (\sigma' \odot e_c) h_j^\top$, and the backward drive of equation 36 is $J_{c,j}^\top e_c = W^\top (\sigma' \odot e_c)$: the child’s error gated by the activation derivative and passed back through the transposed weights. For a multi-layer block, the same backward drive is a vector–Jacobian product through every layer inside it, which must be provided by additional circuitry (Section 7). At

the coarsest extreme the whole network is one node, the update reduces to $\gamma\alpha \nabla_{\theta} f_1$ exactly, and the energy is a wrapper around a full-network backpropagation computation.¹

Accuracy runs the other way. The corrections to the correspondence come from interior nodes, whose equilibria sit $O(\gamma)$ off the forward pass (App. A), so a finer partition has more correction terms. They vanish as $\gamma \rightarrow 0$ at every partition, but granularity sets how small γ must be for a given accuracy.

Measured at two granularities. We compare block granularity (27 nodes) with weight-matrix granularity (74 nodes) on the CIFAR-10 replication checkpoint; the two graphs compute an identical forward pass, and all measurements here are in float64. With a generous relaxation budget ($K_h = 3$, $T_{\text{free}} = 60$, $T_{\text{nudge}} = 30$), $\cos(\Delta\theta_{\text{EP}}, \Delta\theta_{\text{BP}})$ is 1.0000 to four decimals at both granularities and at every γ from 3×10^{-7} to 10^{-4} . At the budget of our EP training run (Table 6), both give 0.9999 at $\gamma = 10^{-4}$ and 1.0000 at $\gamma = 3 \times 10^{-5}$ and 10^{-5} , with update norms within 0.2% of each other; both are about 2% shorter than the backpropagation gradient, a consequence of the budget rather than the partition. The finer partition does raise the correspondence error, as the previous paragraph predicts: the relative velocity deviation is linear in γ at both, with a coefficient 13.7 times larger at the finer one (2.04×10^{-4} against 1.49×10^{-5} at $\gamma = 10^{-4}$). That places the finer partition’s bias at the TF32 floor of Figure 1 rather than below it, so in that arithmetic its operating γ would need to be smaller. At inference both partitions converge in a single sweep in error coordinates ($K_h = 1$ at $dt = 1$, App. E), which costs 7–8% more at the finer one on CPU. These are single-checkpoint measurements; we did not train at weight-matrix granularity, so we cannot say whether its training dynamics differ over many steps.

D.4 THE ADOPTED GRAPH

For the runs of the main text we use block granularity, so that whole modules of the backbone serve as f_k unmodified. The graph has $N = 27$ nodes: twelve encoder blocks chained from the input, the middle block, twelve decoder blocks each taking [previous, skip] as parents with the skip popped from the encoder stack in the architecture’s own order, the output projection, and the ViT head producing o . The time-conditioning pathway inherited from the diffusion backbone is not a node: $V(x)$ is time-independent, so it receives a fixed constant, which each block sees as a trained per-channel bias.²

E ERROR-PARAMETERISED RELAXATION

Rather than relaxing the node states h_k directly, we follow Goemaere et al. (2026) and relax the prediction errors $e_k = f_k(h_{\text{pa}(k)}) - h_k$, recovering states by the topological scan $h_k = f_k(h_{\text{pa}(k)}) - e_k$. This is an invertible change of variables on the same energy, so the two parameterisations share their stationary points. Their curvature, however, is very different, and at the step sizes of a practical simulation only one of them converges in a usable number of sweeps. The error parameterisation is a simulation device rather than dynamics a substrate would run: $\partial E/\partial e$ is a reverse-mode derivative through the scan, and a single sweep from $e = 0$ at $dt = 1$ is one backward pass (Goemaere et al., 2026, Thm. C.9). What it locates, however, is the fixed point of E_{PCN} , which does not depend on how it is found: the local state dynamics equation 36 converge to the same point (below), and every result in this paper is a property of that point.

In the h -parameterisation each state enters both its own error term and the error term of every child that consumes it, so $\partial^2 E/\partial h^2$ has diagonal blocks $I + \sum_c J_{c,j}^\top J_{c,j}$ (plus second-derivative terms in f_c , suppressed by $e_c = O(\gamma)$) and off-diagonal blocks $-J_{c,j}^\top$ for each parent–child pair: its curvature is set by the squared Jacobians of the architecture. In the e -parameterisation the errors are independent coordinates, so $\sum_k \frac{1}{2} \|e_k\|^2$ has Hessian exactly I . The network’s layers now enter only through the injection: h_N is recovered from the errors by the forward scan, so it is a nonlinear

¹With a single node, $h_1 = o$ and $E = \frac{1}{2} \|f_1(x) - h_1\|^2 + \gamma\alpha h_1$. Stationarity in h_1 gives $e_1 = \gamma\alpha$ exactly, and $\partial E/\partial \theta = e_1 \nabla_{\theta} f_1 = \gamma\alpha \nabla_{\theta} f_1$.

²The feedforward reference does the same; injecting the identical constant is what preserves exact agreement with its forward pass. The pathway carries no state and does not relax.

function of every e_k , and the injection $\gamma \alpha h_N$ contributes that function’s curvature, scaled by γ :

$$\frac{\partial^2 E}{\partial e^2} = I + \gamma M, \quad M = \alpha \nabla_e^2 h_N, \quad (38)$$

with M independent of γ .

Measured curvature. On the CIFAR-10 replication checkpoint at the operating $\gamma = 10^{-4}$, power iteration on Hessian–vector products gives a largest Hessian eigenvalue $\mu_{\max}$ of 1.0001 in error coordinates and at least 8×10^4 in state coordinates. Gradient descent at step dt is stable only for $dt < 2/\mu_{\max}$: below 2 in error coordinates, which admits the deadbeat step $dt = 1$, and below 2.4×10^{-5} in state coordinates.

The state parameterisation is impractically slow. At $dt = 1$ the state relaxation diverges within a few sweeps, in float64 as in float32, so the failure is a dynamical instability rather than a precision floor. Below the stability bound it does converge, monotonically and to the same fixed point as the error relaxation, as a shared stationary point requires. But it converges at a measured rate implying a condition number of about 2×10^7 : at $dt = 2 \times 10^{-5}$, reaching equilibrium takes of the order of 10^8 sweeps, against one in error coordinates. This cost is a property of the explicit discretisation rather than of the dynamics: a continuous system settles at the rate of its slowest mode, not its fastest, and Goemaere et al. (2026, App. B.5) make the same point, that the decay which makes discrete state updates slow is an artefact of time discretisation.

Consequences. The error parameterisation is therefore not just a convenience but what makes the relaxation of this network practical. Three further properties follow from its near-identity Hessian, each measured elsewhere in the paper. The relaxation reaches equilibrium in a single deadbeat sweep from feedforward initialisation (App. F). The implicit-gradient solve of Section 3.2 inverts a matrix already close to the identity, so one conjugate-gradient step is as accurate as fifty (App. H). Finally, float32 suffices: at the operating γ the velocity’s deviation from autograd’s sits below the arithmetic’s own floor (Section 4.1). Part of the reason is that the error parameterisation never forms an error as the difference of two nearly equal states, which in state coordinates costs significant digits of precision in every error.

F MEASURING THE VELOCITY CORRESPONDENCE

Measurement. The PCN relaxation velocities v_{PCN} and the FFN autograd velocities v_{FFN} are both evaluated in the stated arithmetic, and each saturates at its own floor. The highest is $\sim 2 \times 10^{-4}$, for float32 tensors whose multiplications run in truncated TF32 precision on the tensor cores — the default A100 math mode, under which all our CIFAR-10 training and FID runs were performed. Forcing those multiplications to full float32 lowers it to $\sim 4 \times 10^{-7}$, and float64 does not saturate anywhere in the measured range, though it requires patching the adopted architecture’s GroupNorm wrapper, which otherwise casts to float32 unconditionally and would impose a spurious float32 floor.

Direction and magnitude. Consider a single interpolant, and let ϕ be the angle between v_{PCN} and v_{FFN} . Resolving v_{PCN} along and perpendicular to v_{FFN} gives components $\|v_{\text{PCN}}\| \cos \phi$ and $\|v_{\text{PCN}}\| \sin \phi$, so the deviation $\Delta v = v_{\text{PCN}} - v_{\text{FFN}}$ has

$$\|\Delta v_{\parallel}\| = \left| \|v_{\text{PCN}}\| \cos \phi - \|v_{\text{FFN}}\| \right|, \quad \|\Delta v_{\perp}\| = \|v_{\text{PCN}}\| \sin \phi. \quad (39)$$

For $\phi \ll 1$ these reduce to $\left| \|v_{\text{PCN}}\| - \|v_{\text{FFN}}\| \right|$ and $\|v_{\text{PCN}}\| \phi$: the parallel component is an error in the velocity’s magnitude, the perpendicular component an error in its direction. Normalising both by $\|v_{\text{FFN}}\|$ gives the relative errors $\epsilon_{\parallel} \approx \left| \|v_{\text{PCN}}\| / \|v_{\text{FFN}}\| - 1 \right|$ and $\epsilon_{\perp} \approx \phi$, the second using $\|v_{\text{PCN}}\| / \|v_{\text{FFN}}\| = 1 + O(\gamma)$. Since the components are orthogonal,

$$\epsilon^2 = \epsilon_{\parallel}^2 + \epsilon_{\perp}^2, \quad (40)$$

where ϵ is the total relative deviation of Figure 1. Figure 2a plots all three: both components are $O(\gamma)$, with direction carrying roughly two-thirds of ϵ^2 across the range. At $\gamma = 10^{-4}$, $\phi \approx 1.1 \times 10^{-5}$ rad, i.e. $\cos(v_{\text{FFN}}, v_{\text{PCN}}) = 1$ to nine decimal places.




Figure 2: Decomposition and budget independence of the velocity deviation, measured as in Figure 1. **(a)** The relative velocity deviation $\epsilon = \|v_{\text{PCN}} - v_{\text{FFN}}\|/\|v_{\text{FFN}}\|$ resolved into components parallel (magnitude error) and perpendicular (direction error) to v_{FFN} . Both scale as $O(\gamma)$, with direction the larger, $\epsilon_{\perp}^2$ carrying roughly two-thirds of ϵ^2 (the curves combine in quadrature). **(b)** At $dt = 1$ the deviation is independent of the number of relaxation sweeps K_h at every γ tested: a single sweep already reaches the equilibrium.

Deadbeat solver. At inference the errors e_k are relaxed for K_h sweeps at $dt = 1$, and Figure 2b shows the deviation is independent of K_h : values in $\{1, 2, 4, 8, 14\}$ agree to within 1%. This follows from three choices acting together. The error parameterisation makes the relaxation Hessian a perturbation of the identity, $\partial^2 E/\partial e^2 = I + \gamma M$ (Appendix E); gradient descent with $dt = 1$ on identity curvature is deadbeat, the iteration map $I - dt H = -\gamma M$ annihilating the error in a single sweep up to $O(\gamma)$; and feedforward initialisation³ starts the system at $e_k = 0$, the exact minimum of the $\gamma = 0$ energy, so only the $O(\gamma)$ displacement the injection causes remains. The first sweep therefore lands within $O(\gamma^2)$ of the equilibrium errors. We read the velocity as $\nabla_x E_{\text{PCN}}$ at fixed e . With the errors held fixed, only the injection term depends on x , so $v_{\text{PCN}} = -\gamma^{-1} \nabla_x E_{\text{PCN}} = -\alpha \nabla_x h_N(x, e)$: the γ cancels, and e enters as the argument of an $O(1)$ function rather than being divided by γ . An $O(\gamma^2)$ error in e therefore moves the velocity by only $O(\gamma^2)$, against the $O(\gamma)$ bias, so the residual deviation in Figure 1 is that bias rather than under-relaxation, as the flatness in K_h confirms. At equilibrium this readout equals the $-\gamma^{-1} J_1^T e_1^*$ of Section 3.1 by the envelope identity; away from equilibrium the two differ. The flatness from $K_h = 1$ is specific to $dt = 1$; smaller steps contract the error by a constant factor each sweep, reaching the same fixed point over a few sweeps. Convergence of the $dt = 1$ iteration requires the eigenvalues of γM to remain below unity, a bound never approached in the measured range. Note though that this is a constraint of the explicit Euler discretisation rather than of the model: continuous-time hardware has no step size, and relaxes by true gradient flow at any γ . All inference measurements in the main text therefore use $K_h = 1$ at $dt = 1$ from feedforward initialisation.

Samples. Figure 3 shows the same sixteen noise realisations sampled by each of the v_{FFN} and v_{PCN} velocity sources; the pairs are visually indistinguishable.

G TRANSFER TO IMAGENET-32

The ImageNet-32 checkpoints use the same architecture and output scale as the CIFAR-10 model, so the PCN construction and the γ convention carry over unmodified; only the trained poten-

³Only the single sweep ($K_h = 1$) depends on this: from unit-Gaussian random initial hidden states — $O(1)$ initial errors, against $O(\gamma)$ at equilibrium — $K_h = 1$ fails to converge, but $K_h \geq 5$ reproduces the feedforward-initialised velocity exactly, unchanged out to $K_h = 60$. The equilibrium is therefore a genuine attractor of the relaxation, not an artifact of the feedforward initialisation.




Figure 3: The same sixteen noise realisations sampled by feedforward autograd velocities (top) and by PCN relaxation velocities (bottom), on our CIFAR-10 phase-2 replication checkpoint at $\gamma = 10^{-4}$ and $\tau_s = 3.25$.




Figure 4: Relative velocity deviation ϵ on the authors’ phase-1 ImageNet-32 checkpoint, measured as in Figure 1. The $O(\gamma)$ law and both arithmetic floors reproduce their CIFAR-10 values, and the clamped potential, with the clamp incorporated into the top prediction function, follows the same law and shares the same floor. At the operating $\gamma = 10^{-6}$ (dotted) the deviation lies far below the floor of the sampling-run arithmetic.

tial $V(x)$ differs. The ImageNet recipe additionally applies a saturating output clamp, sampling with $-\nabla_x c \tanh(V/c)$ for $c = 10^4$. Since the output node carries the potential pre-scale ($V = \alpha o$), we absorb the clamp into the top prediction function as $f'_N = c \tanh(\alpha f_N/c)/\alpha$, so that $\alpha o' = c \tanh(V/c)$ and the injection term $\gamma \alpha o'$ acts on exactly the quantity the feedforward sampler differentiates. The correspondence then applies unchanged, with no additional term.

Figure 4 repeats the sweep of Figure 1 on these checkpoints. The $O(\gamma)$ law holds over the same range, up to $\gamma = 0.1$; the float32 and TF32 floors sit at 3.8×10^{-7} and 2.0×10^{-4} , within ten per cent of their CIFAR-10 values, and the three arithmetics coincide above $\gamma \sim 10^{-2}$ as before. The clamped and unclamped potentials follow the same law. At the operating point $\gamma = 10^{-6}$ the $O(\gamma)$ bias is $\sim 5 \times 10^{-8}$, some four orders of magnitude below the TF32 floor of the sampling runs.




Figure 5: $\cos(\Delta\theta_{\text{EP}}, \Delta\theta_{\text{BP}})$ against the free-phase relaxation budget T_{free} , at six values of γ , in float32 (left) and float64 (right). CIFAR-10 replication checkpoint, 512 flow-matching interpolants, $T_{\text{nudge}} = 6$, $K_h = 1$, $\beta = 30$. Every γ that reaches full fidelity does so at the same budget; the plateau heights, by contrast, are specific to this β (Table 3).

Table 3: Parameter-update fidelity $\cos(\Delta\theta_{\text{EP}}, \Delta\theta_{\text{BP}})$ over nudge strength β (rows) and injection strength γ (columns), varied about the settings of our EP phase-1 training run (Table 6), on the CIFAR-10 replication checkpoint at 512 flow-matching interpolants. Shaded: $\cos < 0.99$. Every failure in the grid belongs to the EP estimator, not the correspondence: the implicit solve, which uses no nudge, is exact to four decimal places up to $\gamma = 10^{-2}$.

	β	10^{-5}	3×10^{-5}	10^{-4}	10^{-3}	10^{-2}	10^{-1}
	0.3	-0.0657	0.2075	0.8423	0.9964	0.9890	0.4395
	1	0.2489	0.8083	0.9807	0.9994	0.9950	0.8746
	3	0.2517	0.9714	0.9970	0.9997	0.9960	0.3620
	10	0.5490	0.9952	0.9996	0.9996	0.9949	0.5965
	30	0.5069	0.9978	0.9999	0.9990	0.9836	0.3374
	100	0.5115	0.9983	0.9997	0.9969	0.9721	0.2578
<i>controls</i>							
	$\Delta\theta_{\text{IFT}}$, float32	1.0000	1.0000	1.0000	1.0000	1.0000	0.9976
	$\Delta\theta_{\text{EP}}$, float64 ($\beta=30$)	1.0000	1.0000	0.9999	0.9982	0.9421	0.6581

H TRAINING FIDELITY: BUDGETS AND NUDGE STRENGTH

The free-phase threshold. Figure 5 plots $\cos(\Delta\theta_{\text{EP}}, \Delta\theta_{\text{BP}})$ against the free-phase budget T_{free} . Where the free phase has not converged the update is not merely degraded but anticorrelated with backpropagation’s, and the recovery is abrupt. The level each curve then plateaus at is a separate matter, set by γ and β together with the arithmetic (Table 3); here we are concerned only with where the transition happens. Across the float32 grid we measured — γ from 5×10^{-5} to 10^{-1} at $\beta = 30$, and $\beta \in \{3, 30, 100\}$ at $\gamma = 10^{-4}$ and 10^{-2} — every cell that reaches full fidelity is below -0.94 at $T_{\text{free}} \leq 6$, above 0.9 at $T_{\text{free}} = 8$, and above 0.99 by $T_{\text{free}} = 10$. Neither γ nor β moves the transition within this range. At $\gamma = 10^{-5}$ the cosine rises but stalls at 0.54 , and below that it never turns positive at any budget; where the update is arithmetic-limited there is no transition to full fidelity to locate. We therefore use $T_{\text{free}} = 14$ throughout, well past the transition, and $T_{\text{nudge}} = 6$, above what the next paragraphs show is needed.

Nudge strength. At the operating $\gamma = 10^{-4}$ the linear nudge holds $\cos \geq 0.99$ from $\beta = 2$ to 10^3 and degrades only to 0.969 at 10^4 ; below $\beta = 2$ it meets the float32 floor, falling from 0.98 at $\beta = 1$ to zero at $\beta = 0.01$, where float64 stays at 1.0000 throughout. The quadratic nudge’s failure above $\beta = 1$ (Figure 6) is a divergence of the nudged relaxation rather than a loss of accuracy: the update




Figure 6: $\cos(\Delta\theta_{EP}, \Delta\theta_{BP})$ against nudge strength at $\gamma = 10^{-4}$, 512 interpolants, $T_{\text{free}} = 14$, $T_{\text{nudge}} = 6$, $K_h = 1$. The quadratic nudge’s relaxation diverges above $\beta = 1$; its values beyond that point are not true gradient estimates.

Table 4: Budget axes that are flat from a single sweep, at $\gamma = 10^{-4}$ on the CIFAR-10 replication checkpoint of Table 1, 512 flow-matching interpolants, otherwise the settings of our EP training run ($K_h = 1$, $T_{\text{free}} = 14$, $T_{\text{nudge}} = 6$, $\beta = 30$) with the named axis varied. Left: conjugate-gradient steps for the implicit solve, $\cos(\Delta\theta_{IFT}, \Delta\theta_{BP})$; centre and right: nudged-phase sweeps and hidden-equilibration steps for EP, $\cos(\Delta\theta_{EP}, \Delta\theta_{BP})$. Dashes mark values not in that axis’s swept grid.

value	n_{CG} (IFT)	T_{nudge} (EP)	K_h (EP)
1	1.0000	0.9998	0.9999
2	1.0000	0.9999	0.9999
3	1.0000	0.9999	0.9999
5	1.0000	—	0.9999
10	1.0000	0.9999	—
20	1.0000	—	—
50	1.0000	—	—

norm leaves the backpropagation reference of 0.323 entirely, reaching 849 at $\beta = 5$. It is a property of the nudge form itself, not of the solver or the arithmetic: reducing the relaxation step size at fixed nudge time makes it worse — at $\beta = 3$ the norm grows by a factor of nearly 400 between $dt = 1$ and 0.1 — and in float64 the collapse occurs at the same β , the two arithmetics lying on top of each other from $\beta = 1$ upward. The linear nudge is unaffected by the step size, and shows no collapse in either arithmetic.

Solver and relaxation budgets. The implicit solve’s parameter update is insensitive to the number of conjugate-gradient steps, giving $\cos(\Delta\theta_{IFT}, \Delta\theta_{BP}) = 1.0000$ at every n_{CG} from 1 to 50 and at every γ tested, which is expected when the relaxation Hessian is a perturbation of the identity (App. E). The truncation therefore contributes nothing, leaving the $O(\gamma)$ correspondence bias as the solve’s only error. The parameter update also remains exact at the small γ where Section 4.1 finds the inference velocity arithmetic-limited, since a floor-level perturbation to v_{PCN} barely rotates the direction of the parameter gradient it produces. The two remaining EP budget axes, T_{nudge} and K_h , are likewise flat from a single sweep (Table 4), for the deadbeat reason of App. E, so T_{free} is the only budget with a threshold, and the cost of an EP update is essentially the free-phase equilibration alone.

I THE HARDWARE-PLAUSIBLE ARCHITECTURE

This appendix specifies the norm- and attention-free architecture of Section 5.1 in enough detail to reimplement, together with the training recipe it needed. It shows only that competitive generation is possible on an architecture an analog substrate could run, not that this particular architecture is the optimal one. We give our reasons for the individual choices where we have them, but which of them are necessary, and at what cost, we leave to future work.

What was removed and what replaced it. The architecture is built from the Energy Matching UNet but with three changes. Firstly, all 51 group-normalisation layers are replaced by the identity. Secondly, all 6 attention blocks are replaced by the identity, which must be done on the built module tree rather than by configuration: the UNet’s middle block hardcodes ResBlock–Attention–ResBlock, so setting the attention resolutions alone cannot produce an attention-free network. Thirdly, the ViT head is replaced by a flat one, which flattens the UNet’s $3 \times 32 \times 32$ output and applies Linear(3072, 512), SiLU, Linear(512, 1). We re-initialise the UNet’s final convolution from $\mathcal{N}(0, 0.02)$ rather than leaving it at the flow-matching codebase’s zero initialisation, after an earlier variant of the head failed to train from zero. We did not test whether the flat head trains without this.

Two further changes are made to the remaining convolutions and residual blocks. Firstly, weight standardisation (Qiao et al., 2020) replaces each convolution’s weights by

$$\hat{w}_i = \frac{g_i w_i}{\sigma_i \sqrt{\text{fan-in} + \epsilon}}, \quad \sigma_i = \text{std}(w_i), \quad (41)$$

per output filter i , with g_i a learned per-filter gain initialised to $\sigma_i \sqrt{\text{fan-in}}$ so that the reparameterisation is exactly the identity at initialisation. No mean subtraction is applied. It is applied to every convolution except the zero-initialised ones — therefore 42 convolutions are standardised and 23 are left plain on the built model. Those 23 are each ResBlock’s second convolution, zero-initialised by the flow-matching codebase so that every block starts as the identity; standardising an all-zero filter would leave both its weight and gain gradients at zero, so it could never leave that initialisation. Secondly, residual scaling replaces each ResBlock’s output by $\text{skip} + \rho(f(x) - \text{skip})$ with $\rho = 1/\sqrt{22} = 0.213$, one over the square root of the number of ResBlocks (De & Smith, 2020). Neither requires computation at inference: $\hat{w}$ depends only on the weights, so a crossbar is programmed with it once and never recomputes a normalisation at runtime, and ρ is a wiring gain.

The resolution-changing path needed no modification: downsampling is a stride-2 convolution rather than pooling, upsampling is nearest-neighbour interpolation followed by a stride-1 convolution, and both are weighted sums like any other convolution.

Module census. The built model contains 42 standardised convolutions, 23 plain convolutions, 22 ResBlocks, 26 linear layers, 69 SiLU activations, and zero group-normalisation, attention or layer-normalisation modules, against 51, 14 and 16 of those respectively in the architecture it replaces. It carries 78.4M parameters against 50.0M, from the flat head’s 3072×512 first linear layer and a base channel width of 192 against 128, which widens every stage through the same multipliers.

How we arrived at this configuration. We reached it by trial rather than by systematic search. Attention and the ViT head could each be removed without much cost, but removing normalisation made training diverge, and weight standardisation was what made the fully stripped network trainable. We did not establish which ingredients are necessary, at what width, or under what schedule; that investigation is left to future work.

Training recipe. This arm is trained by backpropagation, with the settings of Table 6. It differs from the other arms by using a cosine learning rate schedule, annealing to zero at 145k steps, then a constant learning rate of 6×10^{-4} from 145k to 175k steps, after which the second phase is conducted as usual. We arrived at this change by trial after a constant learning rate of 1.2×10^{-3} destabilised partway through the first phase; we did not try a lower constant rate from scratch, and the successful run spent most of its life below 10^{-3} , so the schedule rather than the annealing may be what mattered.

Sampling time. The contrastive phase shifts the optimal sampling time in both architectures, and further in this one. Sampled on a common τ_s ladder after that phase, the full architecture (under all three training rules) is best at $\tau_s = 3.25$ and this one at $\tau_s = 5.0$, and both are much worse at $\tau_s = 1$ than before the contrastive phase. Phase 2 optimises the potential for the Langevin regime at the expense of the noiseless endpoint, and this architecture is affected more.

J EXPERIMENTAL IMPLEMENTATION

This appendix gives the settings needed to reproduce the runs of Section 5, and the implementation choices that the PCN formulation requires but a feedforward one does not. Code is available at <https://anonymous.4open.science/r/PotentialEP>.

Shared recipe. Table 6 gives the settings shared by all arms, and the relaxation settings of the two PCN arms. The three arms of Section 5 share the architecture, optimiser, learning rate, schedule, effective batch size, EMA decay, step budget and data pipeline; the norm- and attention-free arm differs as App. I describes. What they do not share is the random realisation: there is no global seed in the trainer, so each arm has its own initialisation, its own noise samples for the optimal-transport coupling and the contrastive negatives, and its own dropout masks. With one run per arm we have no estimate of run-to-run variance, and the differences in Table 2 should be read with that in mind. The implicit-solve arm holds x fixed at x_t and differentiates the flow loss on $v_{\text{PCN}}(x_t)$ with respect to θ , using the implicit function theorem for the dependence of h^* on θ , so it uses no tether; the tethered objective $\mathcal{L}_{\text{PCN}}$ approaches this fixed- x one as $\gamma/\lambda \rightarrow 0$ (App. B).

Asynchronous relaxation. Nodes are updated in even/odd sweeps (Kerjan et al., 2026), the second half-step seeing the first’s updated values, so one sweep costs two energy evaluations. On a chain this is exactly Gauss–Seidel, since same-parity nodes share no energy term; on the DAG of App. D a skip may connect two same-parity nodes, so the decoupling is only approximate. We adopt it because Kerjan et al. (2026) find synchronous updates substantially worse at the same step size, and sometimes divergent; we did not repeat that comparison here. The fixed point is the same either way.

Frozen dropout. A PCN training step evaluates the network many times: T_{free} sweeps during the free phase and T_{nudge} in each nudged phase, each with K_h inner sweeps, where a feedforward step evaluates it once. Standard dropout would draw a fresh mask at every evaluation, so each sweep would relax a different energy and there would be no single equilibrium for the free phase to reach or for the finite difference to compare. We therefore use one mask per sample, held fixed across every sweep and phase of that sample’s training step. It is drawn as standard dropout draws it, so only the redrawing frequency differs, and for the feedforward arms — a single evaluation per step — the two coincide.

FID protocol. All FIDs are computed from 50k samples against the full training set, with EMA weights, and Euler–Heun integration at $dt = 0.01$. The stopping time τ_s is chosen per model by sampling over a ladder of values and taking the minimum: phase-1 checkpoints are best at the noiseless endpoint $\tau_s = 1$, and after the contrastive phase the optimum moves well past it (App. I). The $\tau_s = 3.25$ chosen as optimal for the backpropagation arm is also near-optimal for both PCN arms, which a paired sweep of τ_s confirms. For the ImageNet-32 checkpoints we use the authors’ $\tau_s = 2.5$.

Measurement resolution and replication. To gauge the measurement’s resolution we evaluated the backpropagation arm’s phase-1 checkpoint at three sampling seeds: the values span 0.064. Changing how the 50k samples are divided across devices moves a fixed-seed evaluation by a similar amount, and unseeded draws of the same checkpoint span 0.15. We therefore treat differences below about 0.2 as unresolved. Our replication reaches 3.54 after phase 2 where the published model reports 3.34; the authors’ released checkpoint scores 3.44 under our protocol, so part of the gap is the sampler and part the replication.

Relaxation step sizes. The two blocks take different steps: the errors are updated at $dt = 1$, deadbeat up to $O(\gamma)$ for the near-identity Hessian $I + \gamma M$ of the error parameterisation (App. E),

Table 5: Phase-1 FID against training step, from 50k samples with EMA weights at $\tau_s = 1.0$. Both arms with a 145k evaluation improve substantially to 100k and are slightly worse by 145k under the constant learning rate, and the EP arm’s 100k checkpoint is better than either arm’s 145k, so its contrastive phase does not start from a weaker model. Repeated evaluations of one checkpoint span up to 0.15, and we treat differences below about 0.2 as unresolved.

Gradient	Phase-1 FID ↓		
	50k steps	100k steps	145k steps
Backpropagation	11.22	6.13	6.42
Implicit solve	12.35	6.45	6.70
EP	12.60	6.23	—

while the visible variables take dt_x , whose block carries the spring and the $O(\gamma)$ curvature of E_{PCN} in x . Phase 1 uses $dt_x = 1$; phase 2 uses $dt_x = 0.5$, paired with $\lambda = 2$ so that $\lambda dt_x = 1$ (below).

The contrastive phase. The contrastive settings are those of Balcerak et al. (2025): 200 Langevin steps at $dt = 0.01$ and $\epsilon_{\text{max}} = 0.01$, half the chains initialised from data and half from noise on the same noise schedule, entering the loss with weight $\lambda_{\text{CD}} = 10^{-3}$ as a trimmed mean over the negatives with the top 10% dropped and the term clamped at 0.02.

For the PCN arms, negatives are drawn by a Langevin chain on a feedforward copy of the network carrying the PCN’s current weights, rather than by relaxation, which our compute budget did not extend to; no parameter gradient passes through that copy, and Section 4.1 shows the two velocity fields agree to a cosine similarity of 1 to nine decimal places. The EP arm’s contrastive phase starts from its 100k phase-1 checkpoint; its 145k checkpoint did not survive the first contrastive updates under any relaxation setting we tried. From 100k, the phase-1 relaxation settings did not carry over either: the free phase stopped converging on some batches after about 60 contrastive steps. The contrastive term raises the flow residual $|g|$, which enlarges the nudge displacement $\beta\gamma|g|/\lambda^2$ of equation 26 and tightens equation 12; the phase-2 settings reduce its prefactor $\beta\gamma/\lambda^2$ by four to compensate. γ falls by $33\times$ and β rises by the same factor, keeping $\beta\gamma$ fixed, and the tether stiffens to $\lambda = 2$, paired with $dt_x = 0.5$ so that $\lambda dt_x = 1$ remains the deadbeat choice of Section 3.2. We did not isolate which of these changes is necessary; in particular, whether reducing β alone by four, at the phase-1 γ , λ and dt_x , would suffice.

K SCOPE OF THE CLAIMS

This appendix states, in one place, what the paper does and does not claim, on four points where a reader might reasonably read more into the results than the evidence supports.

The solver is not the substrate. Every equilibrium in this paper was found by gradient descent in error coordinates (App. E), and a single deadbeat sweep of that solver is algebraically a backward pass (Goemaere et al., 2026, Thm. C.9). That is a choice of simulation method: the claim of Section 3.1 is a property of the equilibrium, not of the route to it. The local state dynamics equation 36 share that equilibrium exactly and converge towards it (App. E), but too slowly under explicit discretisation to run at this scale. What the paper establishes is that a state whose local errors carry $-\nabla_x V$ exists, that the parameter updates EP reads from it match backpropagation’s, and that training on those updates works at CIFAR-10 scale. What it does not establish is that a physical system settles there in useful time; that is the open question a substrate would answer.

What is new. The paper builds on known ingredients: that a PCN’s equilibrium errors implement the backward pass (Whittington & Bogacz, 2017; Millidge et al., 2022), the spring-clamped free phase of GradEP (Gower, 2026), and the linearised nudge of Dual Propagation (Højer et al., 2023). Its contributions are the injection-term construction that turns a scalar-output network into one whose visible layer carries $-\nabla_x V$ (Proposition 1); the joint γ - β window and its derivation (App. C, Table 3); and the demonstration that these compose into competitive generation on CIFAR-10, where prior backpropagation-free generative models remain at MNIST scale.

Table 6: Training settings. The upper block is shared by all arms; the lower gives the relaxation settings of the two PCN arms, with EP’s two phases separately (sweep budgets defined in App. B).

<i>shared</i>			
Optimiser	Adam, $\beta = (0.9, 0.95)$, gradient clipping at norm 1		
Learning rate	1.2×10^{-3} , linear warmup over 10k steps		
Effective batch	1024 (phase 1), 512 (phase 2)		
EMA decay	0.9999 (phase 1), 0.99 (phase 2)		
Steps	145k (phase 1; EP’s contrastive phase starts from its 100k checkpoint) + 2k (phase 2)		
Flow objective	$u = x_1 - x_0$, exact OT per batch, $\sigma = 0$		
Contrastive phase	200 Langevin steps, $dt = 0.01$, $\epsilon_{\max} = 0.01$, $\lambda_{\text{CD}} = 10^{-3}$		
<i>relaxation settings</i>			
	implicit solve	EP (phase 1)	EP (phase 2)
γ	10^{-4}	10^{-4}	3×10^{-6}
λ	—	1	2
β	—	30	10^3
K_h	2	1	1
T_{free}	—	14	14
T_{nudge}	—	6	6
nudge	—	linear	linear
dt_x	—	1	0.5
dt (errors)	1	1	1
n_{CG}	10	—	—

Where a backward pass remains in training. The EP arm’s parameter updates in both phases come from measurements at equilibrium. Its phase-2 negatives, however, were sampled by a feedforward twin of the network using autograd velocities, since relaxing 200 Langevin steps per negative was beyond our budget. No parameter gradient passes through the twin, and the two velocity fields agree to a cosine similarity of 1 to nine decimal places (Section 4.1), so we expect replacing it with relaxation changes the negatives by an amount below the sampler’s own noise, though we did not test this. “Without backpropagation” in Section 5 refers to how the weights were updated, not to every computation the training loop performed.

The hardware-plausible architecture is an existence proof. The norm- and attention-free network of Section 5.1 is wider than the reference, used a different schedule, and was trained and sampled feedforward only. Its FID of 3.51 shows that neither component is needed for competitive generation; it does not show that they can be removed at matched cost, nor that the resulting network trains under EP. Both are left to future work, as is the run-to-run variance of every arm in Table 2, each of which is a single run.

L GENERATED SAMPLES

Figures 7–10 show 128 samples from each of the four trained models of Section 5, drawn from each model’s phase-2 checkpoint under its own native inference — feedforward autograd for the two backpropagation-trained models, PCN relaxation for the other two. All four share a sampling seed, so the panel in a given position starts from the same x_0 in every figure; the sampling time is each model’s own optimum and is given per figure.




Figure 7: CIFAR-10 samples from the backpropagation arm, FID 3.54 at $\tau_s = 3.25$.




Figure 8: CIFAR-10 samples from the implicit-solve arm, FID 3.64 at $\tau_s = 3.25$. Parameter gradients come from implicit differentiation of the PCN’s equilibrium rather than from backpropagation through a forward pass.




Figure 9: CIFAR-10 samples from the EP arm, FID 3.64 at $\tau_s = 3.25$ — a model trained without backpropagation in either phase.




Figure 10: CIFAR-10 samples from the norm- and attention-free architecture of Section 5.1, FID 3.51 at $\tau_s = 5.0$. Every operation in this network is a fixed weighted sum or a pointwise nonlinearity.